

# 深度学习：基础，前沿，应用，评价

吴伟

lazyparser@gmail.com  
中国科学院软件研究所

# 关于我

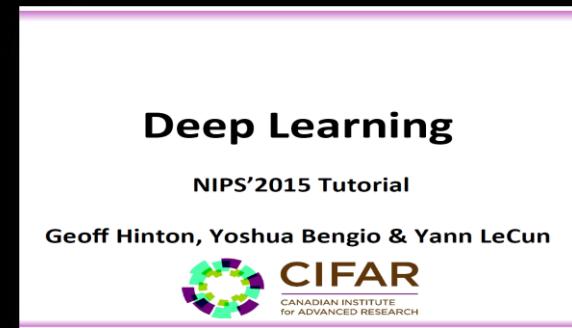
软件研究所 | 神经芯片 | 即时编译 | 机器人 | 人脸识别

开源贡献者 | hellogcc | Mozilla(s) | mxnet | TK1



基础 | 前沿 | 应用 | 评价

先缩小一下入门的范围



**Deep Learning: machine learning algorithms based on learning multiple levels of representation / abstraction.**

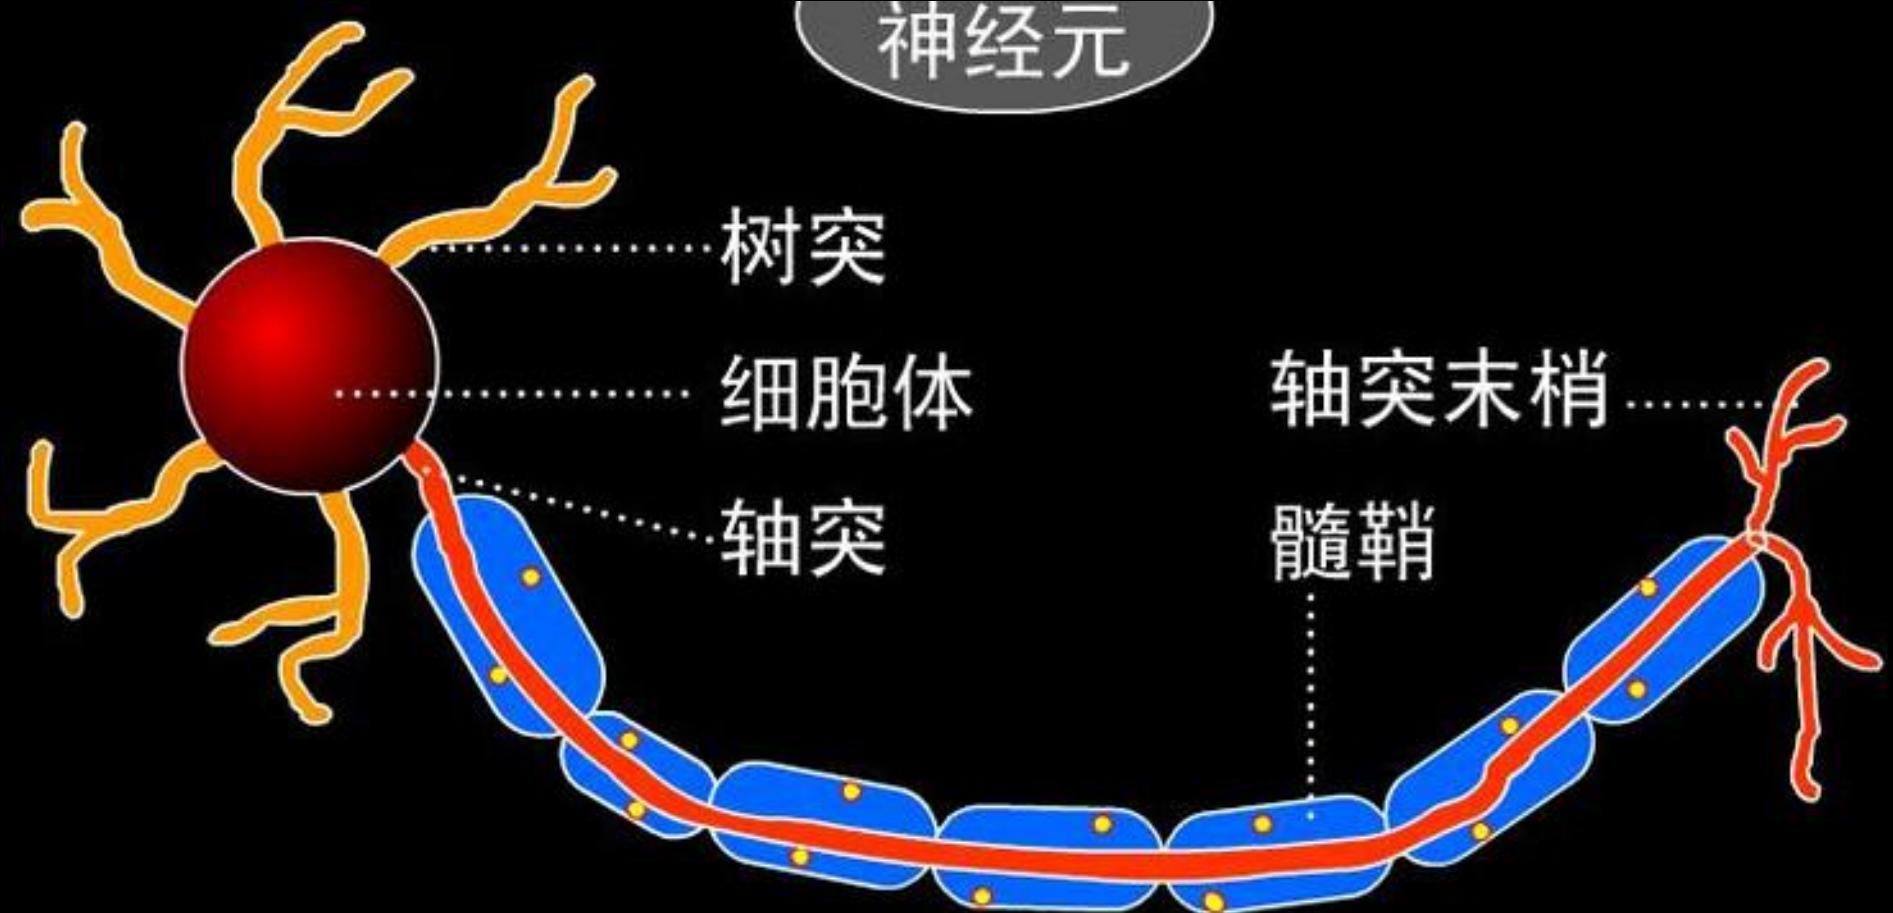
先缩小一下入门的范围

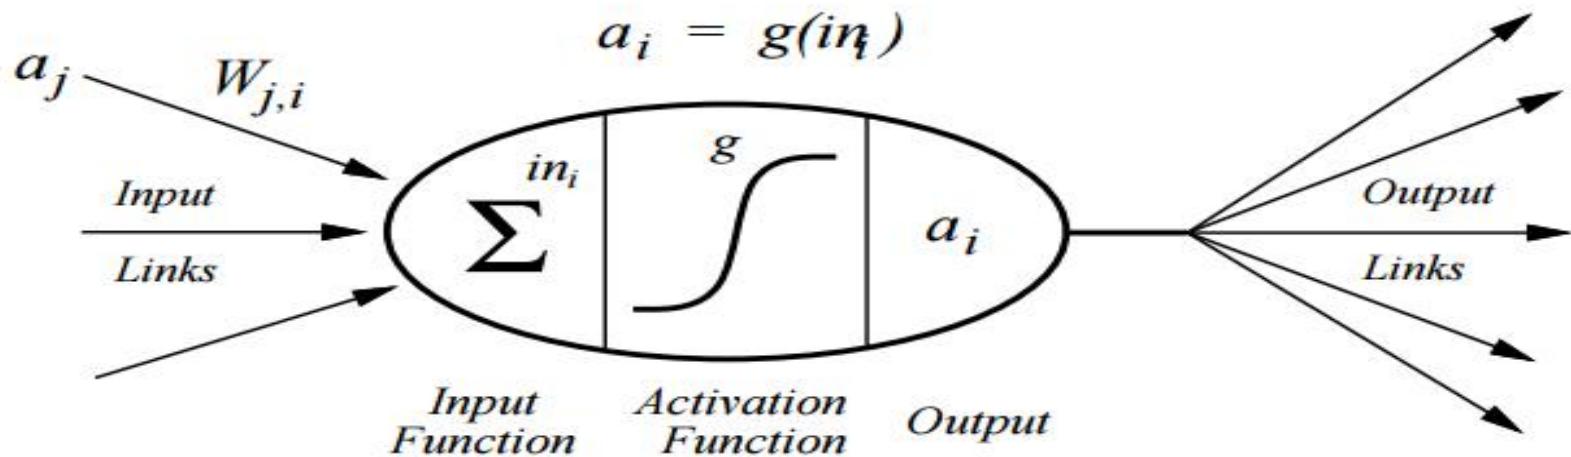
深度学习：深度神经网络

先缩小一下入门的范围

深度学习 : 深度**神经**网络

# 神经元





$$a_i = g\left(\sum_j W_{j,i} a_j\right)$$

$$p(\mathbf{v}) = \sum_{\mathbf{h}} \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{u}, \mathbf{g}} e^{-E(\mathbf{u}, \mathbf{g})}}. \quad (1)$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i a_i v_i - \sum_j b_j h_j - \sum_{i,j} w_{ij} v_i h_j, \quad (2)$$

$$\begin{aligned} p(\mathbf{v}|\mathbf{h}) &= \prod_i p(v_i|\mathbf{h}) \quad \text{and} \quad p(v_i = 1|\mathbf{h}) = \text{sigm} \left( a_j + \sum_j h_j w_{ij} \right), \\ p(\mathbf{h}|\mathbf{v}) &= \prod_j p(h_j|\mathbf{v}) \quad \text{and} \quad p(h_j = 1|\mathbf{v}) = \text{sigm} \left( b_j + \sum_i v_i w_{ij} \right), \end{aligned} \quad (3)$$

$$\begin{aligned} p(v_i = x|\mathbf{h}) &= \frac{1}{\sigma_i \sqrt{2\pi}} \cdot e^{-\frac{(x - a_i - \sigma_i \sum_j w_{ij} h_j)^2}{2\sigma_i^2}}, \\ p(h_j = 1|\mathbf{v}) &= \text{sigm} \left( b_j + \sum_i \frac{v_i}{\sigma_i} w_{ij} \right). \end{aligned} \quad (4)$$

$$p(\mathbf{v}) = \sum_{\mathbf{h}} \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{u}, \mathbf{g}} e^{-E(\mathbf{u}, \mathbf{g})}}. \quad (1)$$

~~$$E(\mathbf{v}, \mathbf{h}) = - \sum_i a_i v_i - \sum_j b_j h_j - \sum_{i,j} w_{ij} v_i h_j, \quad (2)$$~~

~~$$\begin{aligned} p(\mathbf{v}|\mathbf{h}) &= \prod_i p(v_i|\mathbf{h}) \quad \text{and} \quad p(v_i = 1|\mathbf{h}) = \text{sigm} \left( a_j + \sum_j h_j w_{ij} \right), \\ p(\mathbf{h}|\mathbf{v}) &= \prod_j p(h_j|\mathbf{v}) \quad \text{and} \quad p(h_j = 1|\mathbf{v}) = \text{sigm} \left( b_j + \sum_i v_i w_{ij} \right), \end{aligned} \quad (3)$$~~

~~$$\begin{aligned} p(v_i = x|\mathbf{h}) &= \frac{1}{\sigma_i \sqrt{2\pi}} \cdot e^{-\frac{(x - a_i - \sigma_i \sum_j w_{ij} h_j)^2}{2\sigma_i^2}}, \\ p(h_j = 1|\mathbf{v}) &= \text{sigm} \left( b_j + \sum_i \frac{v_i}{\sigma_i} w_{ij} \right). \end{aligned} \quad (4)$$~~

# 现场编一个故事背景

某公司老板赵总，有一天突然有了一个灵感：

设计开发一个**人工智能**软件，利用当天股价的涨跌来预测第二天房价的涨跌。

# 现场编一个故事背景

某公司老板赵总，有一天突然有了一个灵感：  
设计开发一个人工智能软件，利用当天股价的  
涨跌来预测第二天房价的涨跌。

于是找来工程师小李，要求实现这个app：

if 股价涨 then 显示房价涨；  
else 显示房价不会涨；

# 产品原型

```
if (x > 0)
    y = 1
else /* x <= 0 */
    y = 0
```

产品原型

$$f(x) = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

if (x > 0)

y = 1

else /\* x <= 0 \*/

y = 0

# 产品原型（高级版）

```
if (x > 0)
```

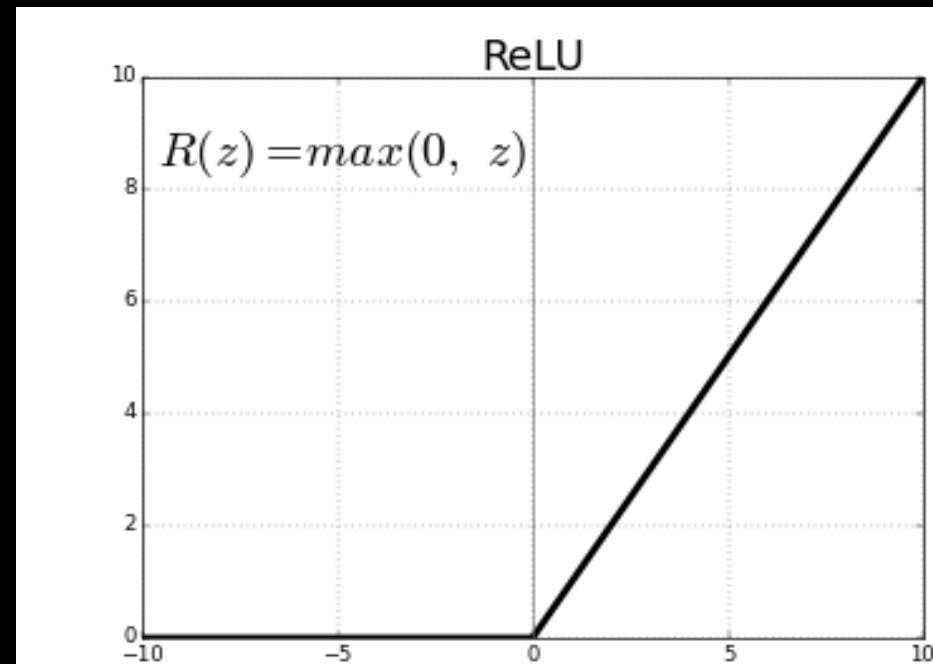
```
    y = x
```

```
else /* x <= 0 */
```

```
    y = 0
```

# 产品原型（高级版）

```
if (x > 0)  
    y = x  
else /* x  
    y = 0
```



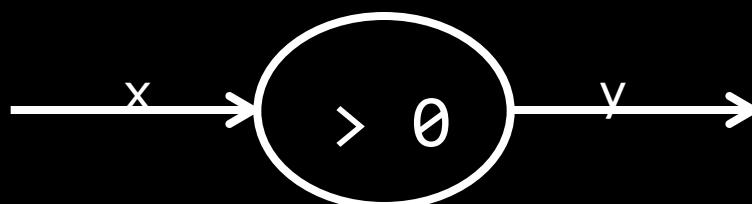
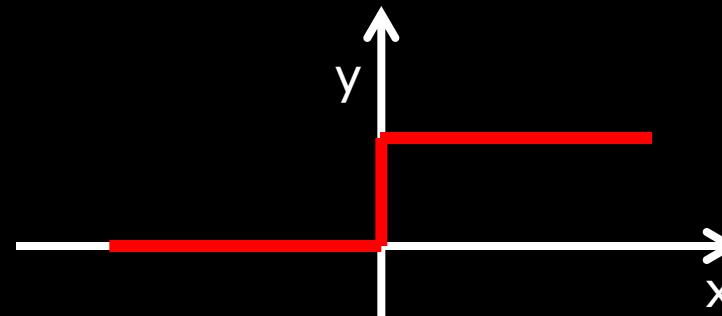
# 产品原型

```
if (x > 0)
```

```
    y = 1
```

```
else
```

```
    y = 0
```



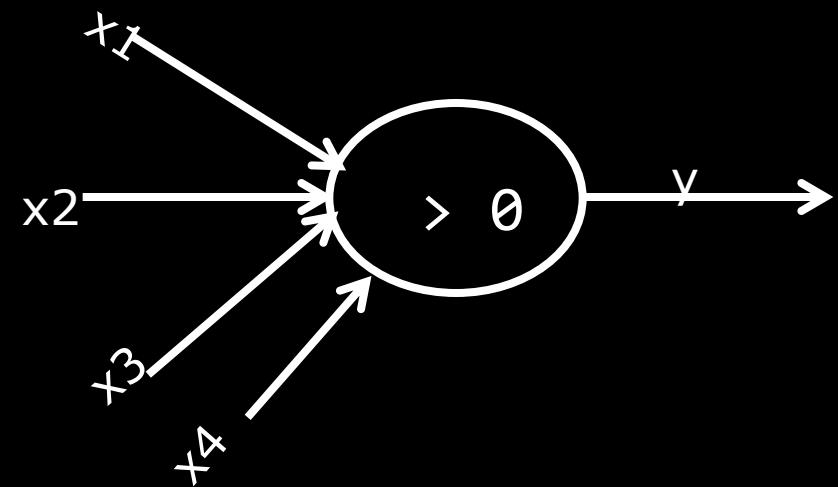
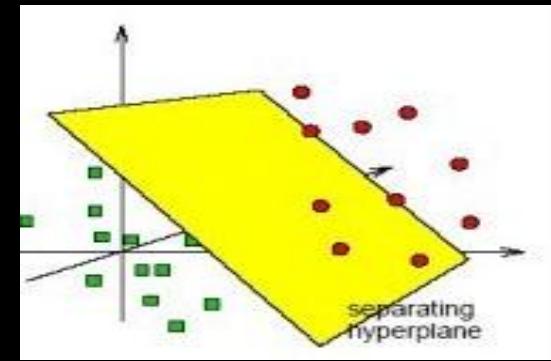
# 故事背景

小李很快做出来了产品的原型，但是显然，  
这个软件是预测不准的.....

赵总思考了一下，觉得应该把油价、汇率、金价  
都考虑进去，这样预测就准了。

# 产品原型

```
foo(x1,x2,x3,x4) {  
    return x1+x2+x3+x4;  
}  
if (foo(x) > 0)  
    y = 1  
else  
    y = 0
```



# 故事背景

小李很快迭代出了产品，但是显然，  
这个软件还是预测不准的.....

赵总思考了一下，觉得股价、油价、汇率、金价  
对房价的影响是不一样的，也就是权重 ( weights )  
是不同的。

# 故事背景

小李明白权重的意思，但不知道具体怎么取值。

赵总根据经验，告诉小李：

股价、油价、汇率、金价，权重分别是

1.0, 0.1, 0.2, 0.5

# 产品原型

```
foo(x1, x2, x3, x4) {
```

```
    return 1.0*x1 + 0.1*x2  
        + 0.2*x3 + 0.5*x4;
```

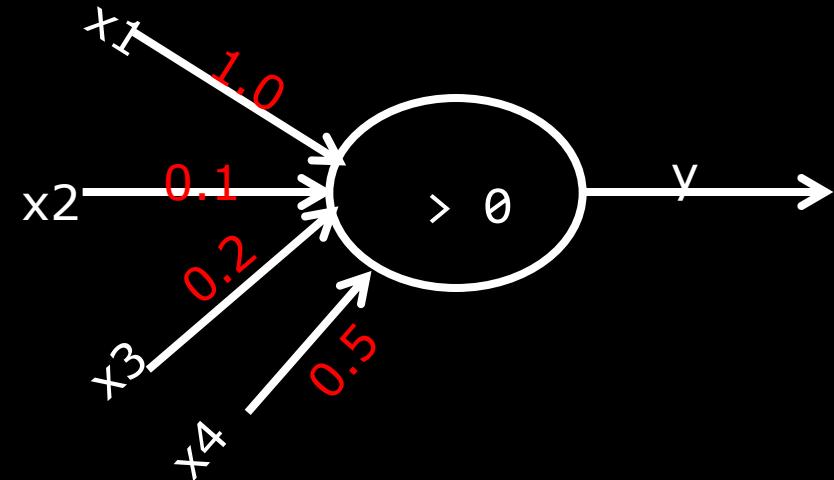
```
}
```

```
if (foo(x) > 0)
```

```
    y = 1
```

```
else
```

```
    y = 0
```



## 故事背景

赵总进一步觉得，长远来看，就算市场基本面不变，房价也在缓慢上涨。所以只有指标看跌到一定程度，才能预测说是房价要跌。

赵总根据经验，告诉小李：

只有指标达到了-0.5了，才能算是要跌。

# 产品原型

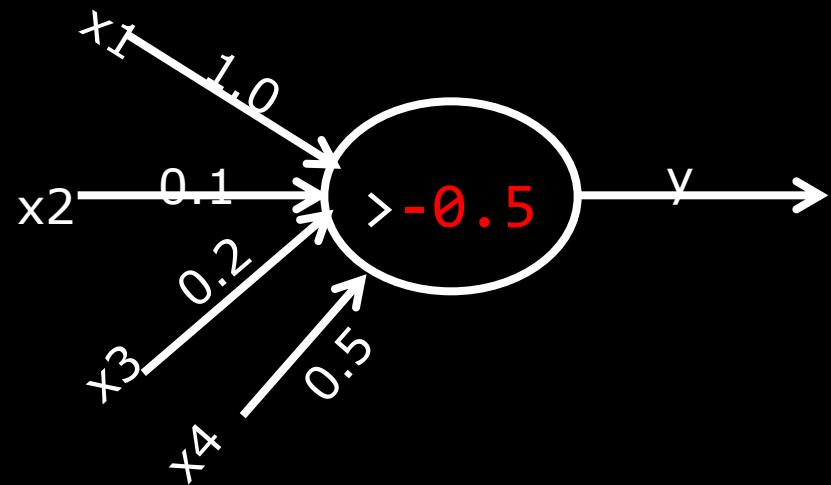
```
foo(x1, x2, x3, x4) {  
    return 1.0*x1 + 0.1*x2  
        + 0.2*x3 + 0.5*x4;  
}
```

```
if (foo(x) > -0.5)
```

```
    y = 1
```

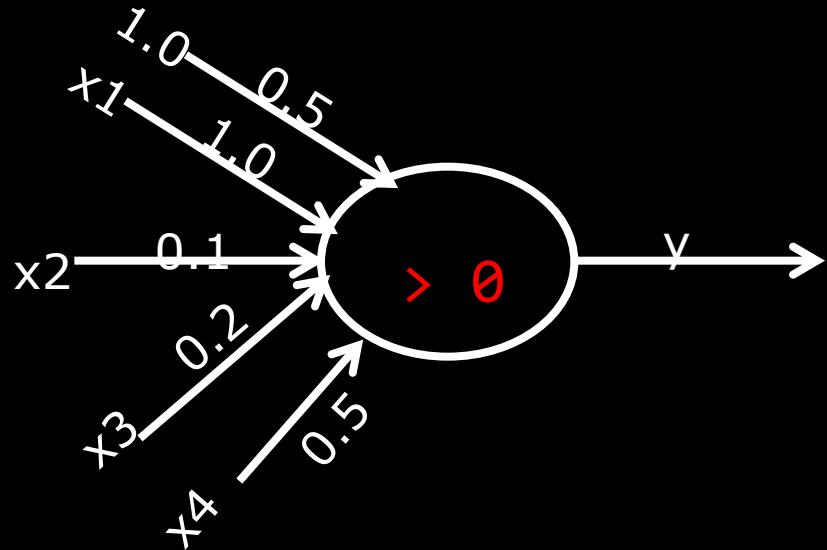
```
else
```

```
    y = 0
```



# 产品原型-小李重构了一下

```
foo(x1, x2, x3, x4) {  
    return 0.5*1 + 1.0*x1 + 0.1*x2  
        + 0.2*x3 + 0.5*x4;  
}  
  
if (foo(x) > 0)  
    y = 1  
else  
    y = 0
```



# 产品原型-小李重构了一下

```
w = [+0.5, 1.0, 0.1, 0.2, 0.5]
```

```
x = [ 1, x1, x2, x3, x4]
```

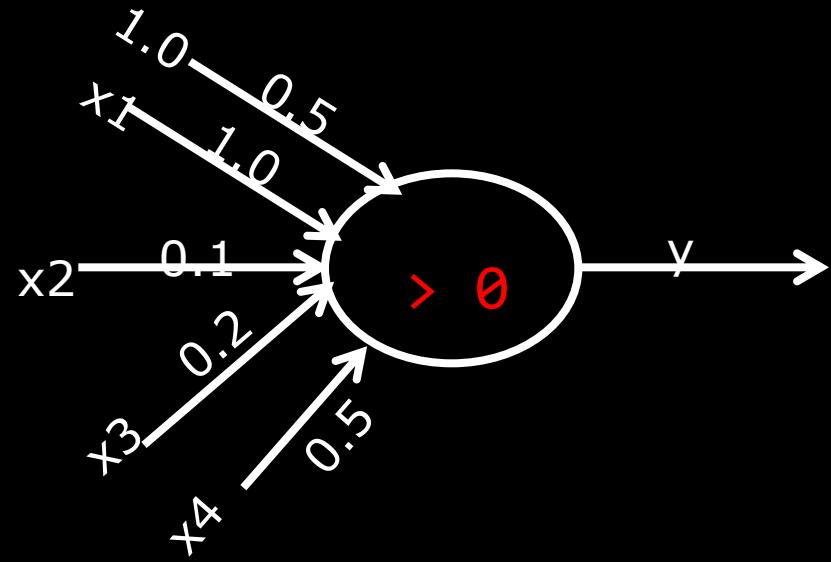
```
foo(x) { return sum(dot_mult(w,x)) }
```

```
if (foo(x) > 0)
```

```
    y = 1
```

```
else
```

```
    y = 0
```



# 产品原型-小李重构了一下

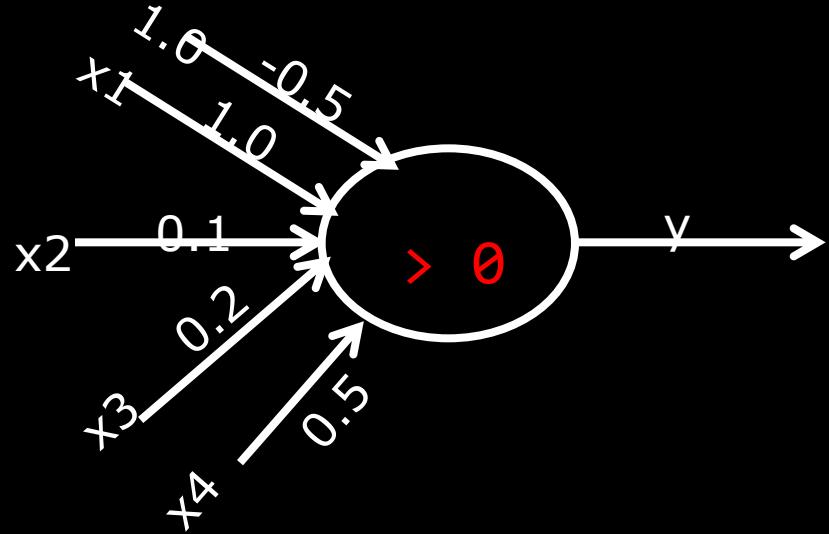
```
w = [+0.5, 1.0, 0.1, 0.2, 0.5]
```

```
x = [ 1, x1, x2, x3, x4]
```

```
foo(x) { return sum(dot_mult(w,x)) }
```

```
if (foo(x) > 0)
    y = 1
else
    y = 0
```

激活函数  
Activation function

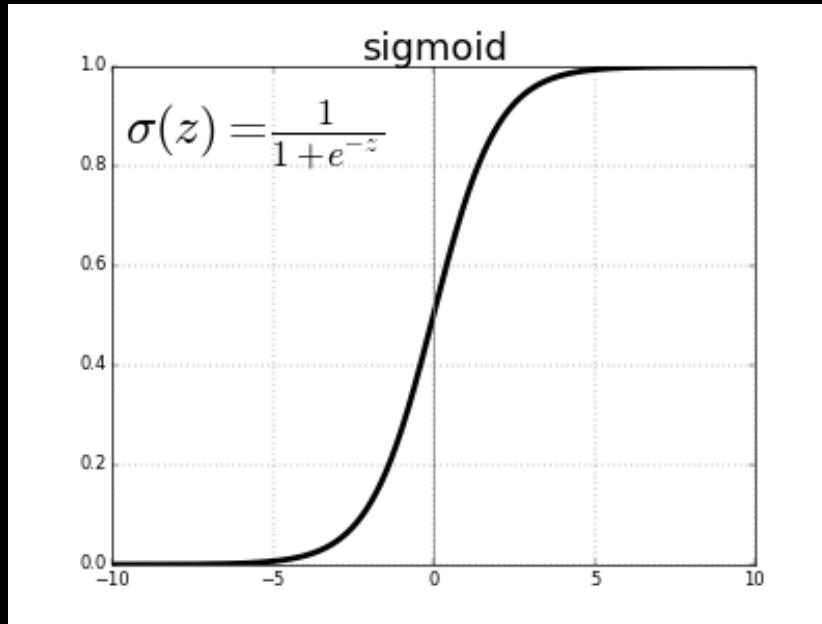


# 故事背景

软件迭代了四个版本了，测起来还是不准。

赵总读了微信公众号的深度学习文章，告诉小李：  
这个关系很复杂，我们之前用的是线性关系，  
不能有效的表示这种复杂的关系，  
我们要用“非线性关系”，用最流行的Sigmoid。

# Sigmoid 函数（两种解释，类比HTML5）



神经网络用Sigmoid函数，其实是为了后面反向传播算法的时候计算微分方便

# 产品原型-小李重构了一下

```
w = [+0.5, 1.0, 0.1, 0.2, 0.5]  
x = [ 1, x1, x2, x3, x4]
```

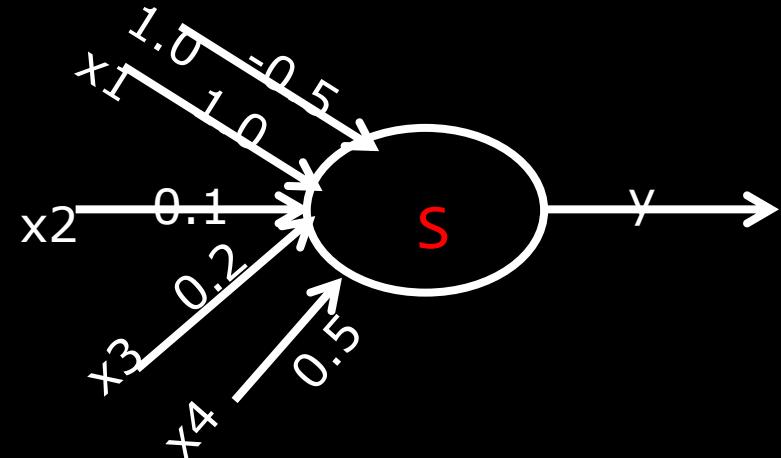
```
foo(x) { return sum(dot_mult(w,x)) }
```

```
if (sigmoid(foo(x))>0.5)
```

```
    y = 1
```

```
else
```

```
    y = 0
```



# 激活 ( Activation ) 函数大家族

Name	Plot	Equation	Derivative (with respect to $x$ )	Range	Order of continuity	Monotonic	Derivative Monotonic	Approximates identity near the origin
Identity		$f(x) = x$	$f'(x) = 1$	$(-\infty, \infty)$	$C^\infty$	Yes	Yes	Yes
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$	$\{0, 1\}$	$C^{-1}$	Yes	No	No
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$	$(0, 1)$	$C^\infty$	Yes	No	No
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$	$(-1, 1)$	$C^\infty$	Yes	No	Yes
Arctan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$	$(-\frac{\pi}{2}, \frac{\pi}{2})$	$C^\infty$	Yes	No	Yes
Softsign [7][8]		$f(x) = \frac{x}{1 +  x }$	$f'(x) = \frac{1}{(1 +  x )^2}$	$(-1, 1)$	$C^1$	Yes	No	Yes
Rectifier (ReLU) <sup>[9]</sup>		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$[0, \infty)$	$C^0$	Yes	Yes	No
Parameteric Rectified Linear Unit (PReLU) <sup>[10]</sup>		$f(\alpha, x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(\alpha, x) = \begin{cases} \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\infty, \infty)$	$C^0$	Yes iff $\alpha \geq 0$	Yes	Yes iff $\alpha = 1$
Exponential Linear Unit (ELU) <sup>[11]</sup>		$f(\alpha, x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(\alpha, x) = \begin{cases} f(x) + \alpha & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\alpha, \infty)$	$C^1$ when $\alpha = 1$	Yes iff $\alpha \geq 0$	Yes iff $0 \leq \alpha \leq 1$	Yes iff $\alpha = 1$
SoftPlus <sup>[12]</sup>		$f(x) = \ln(1 + e^x)$	$f'(x) = \frac{1}{1 + e^{-x}}$	$(0, \infty)$	$C^\infty$	Yes	Yes	No
Bent identity		$f(x) = \frac{\sqrt{x^2 + 1} - 1}{2} + x$	$f'(x) = \frac{x}{2\sqrt{x^2 + 1}} + 1$	$(-\infty, \infty)$	$C^\infty$	Yes	Yes	Yes
SoftExponential <sup>[13]</sup>		$f(\alpha, x) = \begin{cases} -\frac{\log_e(1 - \alpha(x + \alpha))}{\alpha} & \text{for } \alpha < 0 \\ x & \text{for } \alpha = 0 \\ \frac{e^{\alpha x} - 1}{\alpha} + \alpha & \text{for } \alpha > 0 \end{cases}$	$f'(\alpha, x) = \begin{cases} \frac{1}{1 - \alpha(x + \alpha)} & \text{for } \alpha < 0 \\ e^{\alpha x} & \text{for } \alpha \geq 0 \end{cases}$	$(-\infty, \infty)$	$C^\infty$	Yes	Yes	Yes iff $\alpha = 0$
Sinusoid		$f(x) = \sin(x)$	$f'(x) = \cos(x)$	$[-1, 1]$	$C^\infty$	No	No	Yes

# 故事背景

软件迭代了五个版本了，测起来还是不准。

赵总觉得之所以测不准，还是因为模型太简单了。

现在深度学习，动辄就说“一百多层”，看看自己，一层都没有，肯定不行。

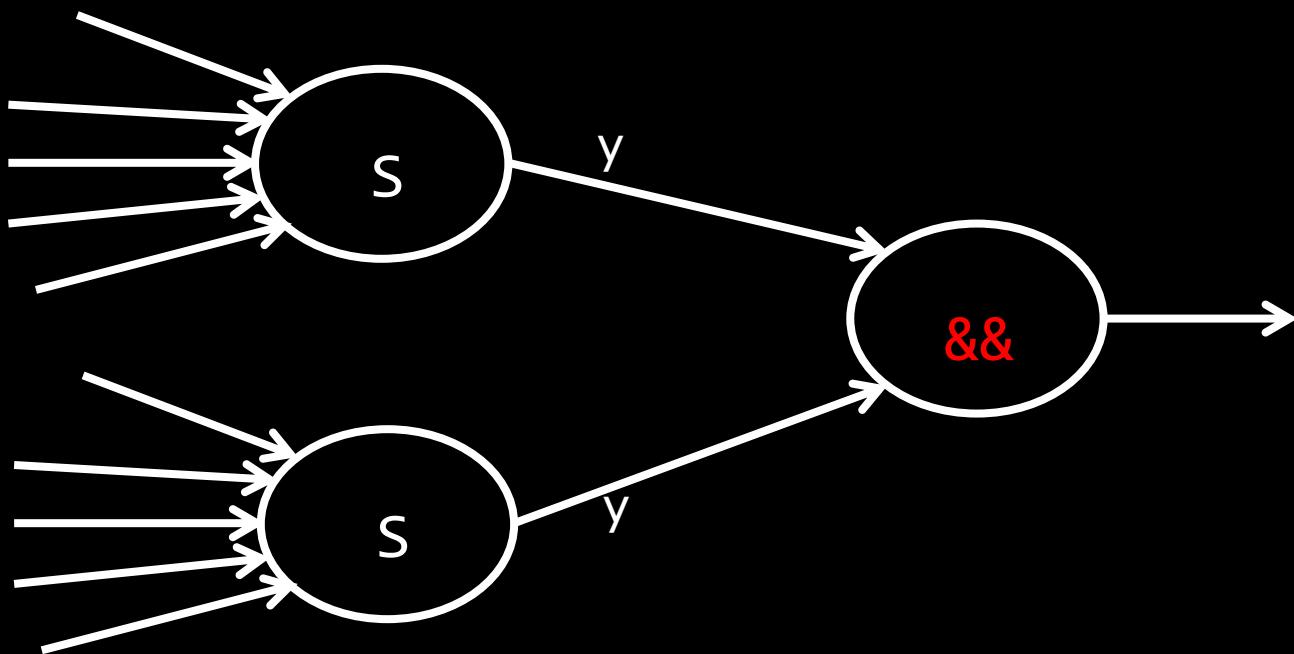
# 故事背景

一天早晨，赵总突然想到，买房对谁都是大事，  
房价预测岂能用一个公式“一刀切”？

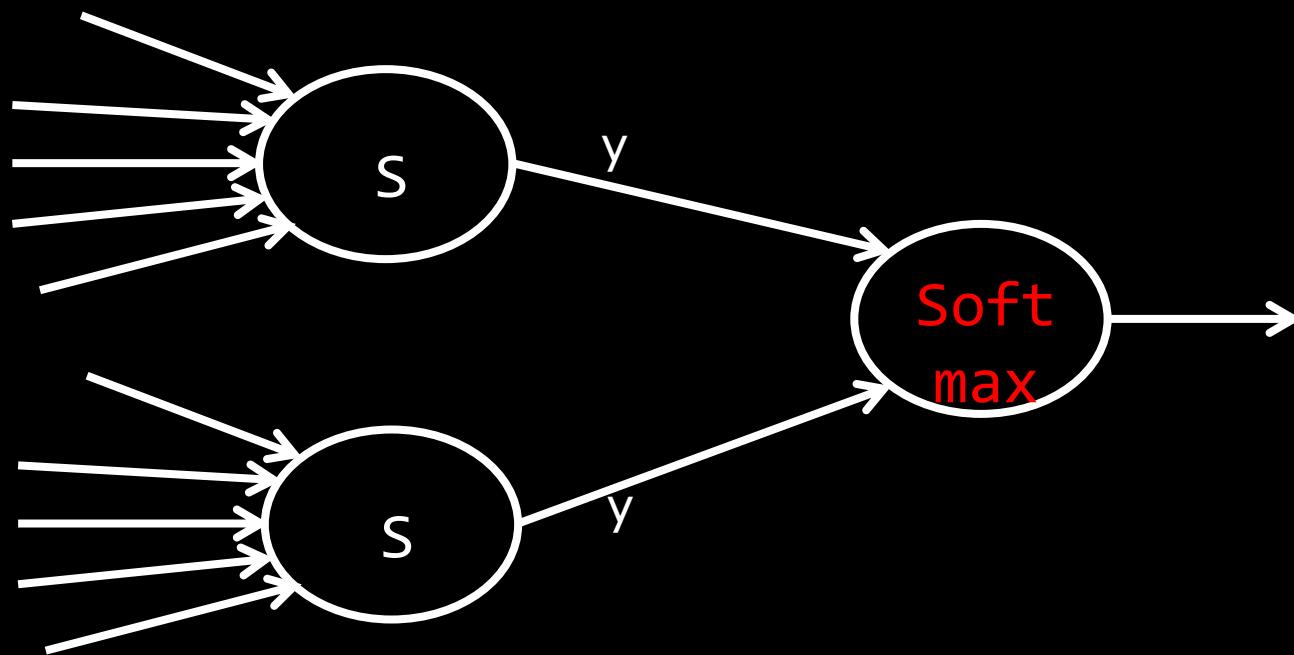
于是把小李找来，告诉他：

我们要进一步完善，用多个不同的计算公式，  
综合最终的结果给出预测结果。

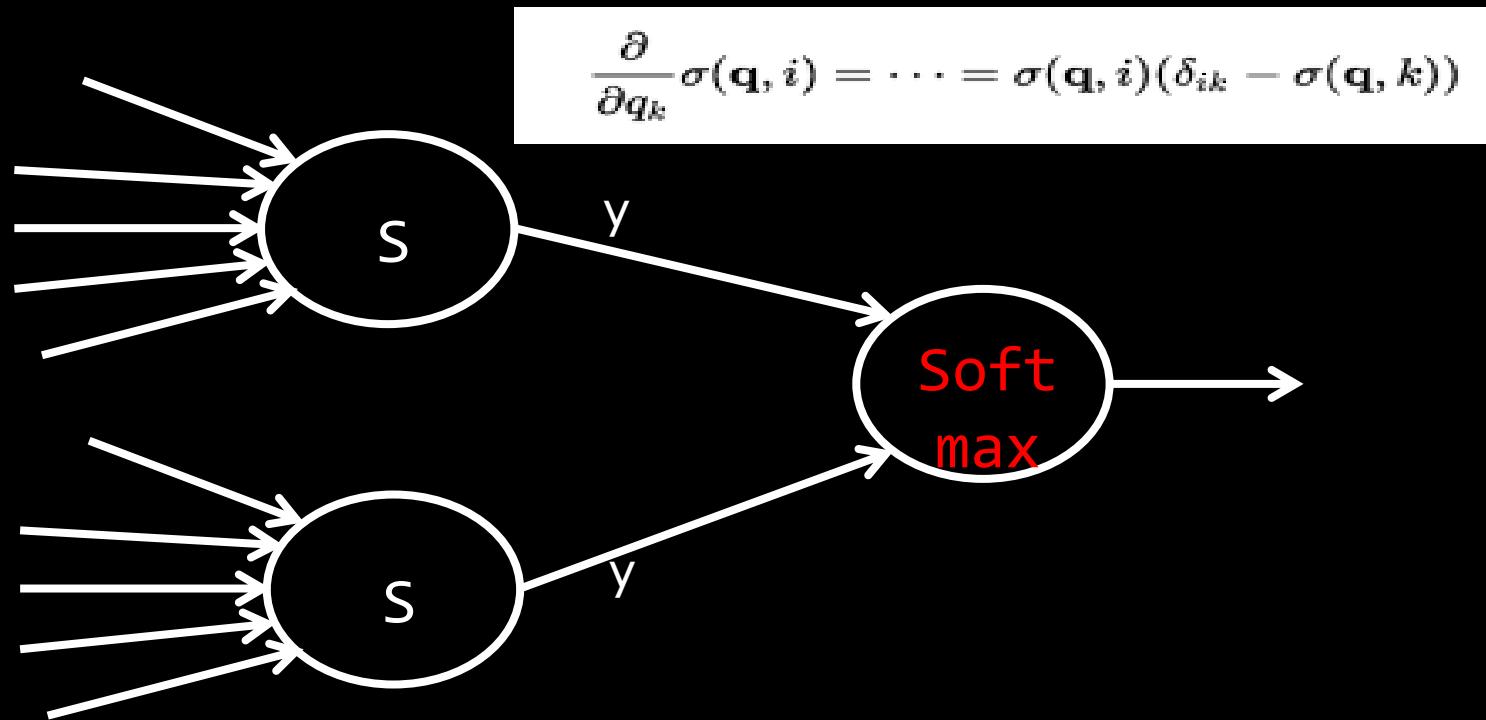
小李很快复制粘贴出来了结果，全说涨才涨



照猫画虎将最后的投票函数改成了 softmax



# 照猫画虎将最后的投票函数改成了 softmax



# 故事背景

有神经元，有层，是神经网络了！但是.....

两个输入的神经元，权重一模一样.....

赵总将这个问题交给了小李，让小李自己去完成。

# 故事背景

现在小李手里有神经网络（**模型**），有房价、股价、油价、汇率和金价的历史数据（**训练集**），也有赵总给的经验权重（**初始值**）。

小李**编**了几个权重，效果都不好。接下来怎么办，小李也不知道了……

# 大神出场：反向传播算法&BP神经网络

## Learning representations by back-propagating errors

David E. Rumelhart\*, Geoffrey E. Hinton†  
& Ronald J. Williams\*

\* Institute for Cognitive Science, C-015, University of California,  
San Diego, La Jolla, California 92093, USA

† Department of Computer Science, Carnegie-Mellon University,  
Pittsburgh, Philadelphia 15213, USA

---

We describe a new learning procedure, back-propagation, for networks of neurone-like units. The procedure repeatedly adjusts the weights of the connections in the network so as to minimize a measure of the difference between the actual output vector of the net and the desired output vector. As a result of the weight adjustments, internal 'hidden' units which are not part of the input or output come to represent important features of the task domain, and the regularities in the task are captured by the interactions of these units. The ability to create useful new features distinguishes back-propagation from earlier, simpler methods such as the perceptron-convergence procedure<sup>1</sup>.

more difficult when we introduce hidden units whose actual or desired states are not specified by the task. (In perceptrons, there are 'feature analysers' between the input and output that are not true hidden units because their input connections are fixed by hand, so their states are completely determined by the input vector: they do not learn representations.) The learning procedure must decide under what circumstances the hidden units should be active in order to help achieve the desired input-output behaviour. This amounts to deciding what these units should represent. We demonstrate that a general purpose and relatively simple procedure is powerful enough to construct appropriate internal representations.

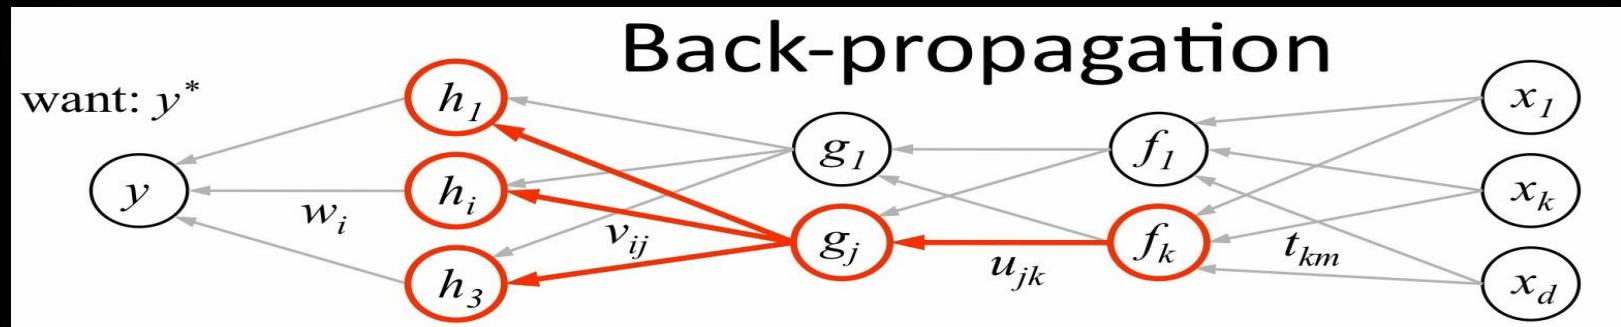
The simplest form of the learning procedure is for layered networks which have a layer of input units at the bottom; any number of intermediate layers; and a layer of output units at the top. Connections within a layer or from higher to lower layers are forbidden, but connections can skip intermediate layers. An input vector is presented to the network by setting the states of the input units. Then the states of the units in each layer are determined by applying equations (1) and (2) to the connections coming from lower layers. All units within a layer have their states set in parallel, but different layers have their states set sequentially, starting at the bottom and working upwards until the states of the output units are determined.

The first two authors contributed equally to this work.

# 思路：用历史数据来修正权重

11.1	11.2	11.3	11.4	.....	11.30
股价	涨	跌	涨	.....	涨
汇率	跌	涨	跌	.....	跌
金价	涨	跌	涨	.....	涨
油价	跌	涨	跌	.....	跌
房价 (有监督)	涨	涨	涨	涨	涨

( 不理解没关系，该算法已经不用自己实现了 )



<https://www.youtube.com/watch?v=An5z8IR8asY>

<https://i.ytimg.com/vi/An5z8IR8asY/maxresdefault.jpg>

# 进一步学习：吴恩达在Coursera上的机器学习课程

## Backpropagation algorithm

> Training set  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$

Set  $\Delta_{ij}^{(l)} = 0$  (for all  $l, i, j$ ).

(use to compute  $\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta)$ )

For  $i = 1$  to  $m$   $\leftarrow (x^{(i)}, y^{(i)})$

Set  $a^{(1)} = x^{(i)}$

Perform forward propagation to compute  $a^{(l)}$  for  $l = 2, 3, \dots, L$

Using  $y^{(i)}$ , compute  $\delta^{(L)} = a^{(L)} - y^{(i)}$

Compute  $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$

$\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + a_j^{(l)} \delta_i^{(l+1)}$

$\Delta_{ij}^{(l)} := \Delta_{ij}^{(l)} + \delta^{(l+1)} (a^{(l)})^T$

$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)} + \lambda \Theta_{ij}^{(l)}$  if  $j \neq 0$

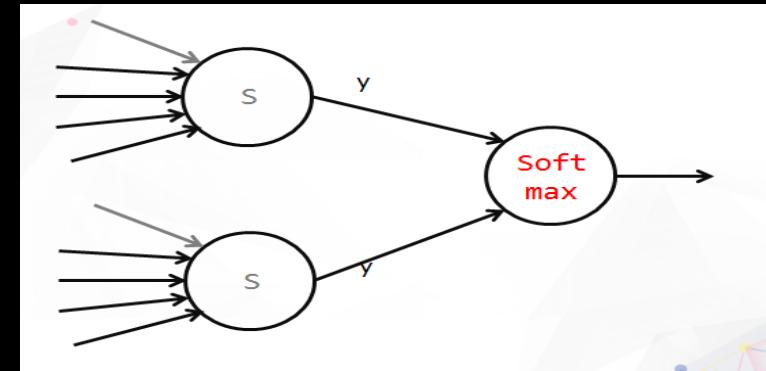
$\frac{\partial}{\partial \Theta_{ij}^{(l)}} J(\Theta) = D_{ij}^{(l)}$

$D_{ij}^{(l)} := \frac{1}{m} \Delta_{ij}^{(l)}$  if  $j = 0$

# 故事背景

小李激动的训练完，效果确实比以前好了很多，  
但是.....

权重还是都是一样的嘛！



不起眼但是很重要：随机初始化

# 随机化在训练中经常需要被使用到

**Backprop in Practice** Y LeCun

- Use ReLU non-linearities
- Use cross-entropy loss for classification
- Use Stochastic Gradient Descent on minibatches
- Shuffle the training samples (← very important)
- Normalize the input variables (zero mean, unit variance)
- Schedule to decrease the learning rate
- Use a bit of L1 or L2 regularization on the weights (or a combination)
  - ▶ But it's best to turn it on after a couple of epochs
- Use "dropout" for regularization
- Lots more in [LeCun et al. "Efficient Backprop" 1998]
- Lots, lots more in "Neural Networks, Tricks of the Trade" (2012 edition) edited by G. Montavon, G. B. Orr, and K-R Müller (Springer)
- More recent: Deep Learning (MIT Press book in preparation)

# 故事背景

小李用了反向传播算法，准确率上去了。

老板很开心，让他加层，并问小李里面权重的意思。

小李表示不知道，于是去问各路大神.....



“不知道”



“很复杂”

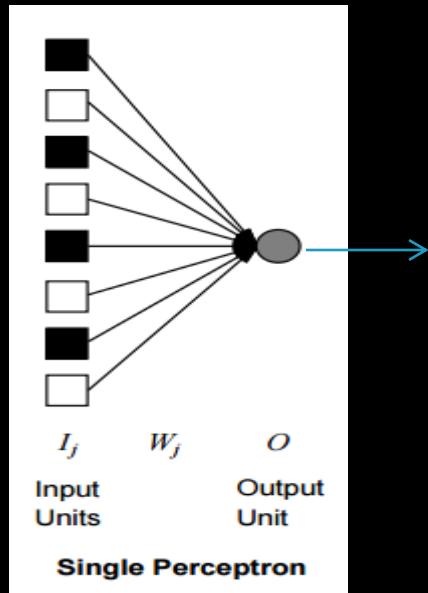


“不好说”

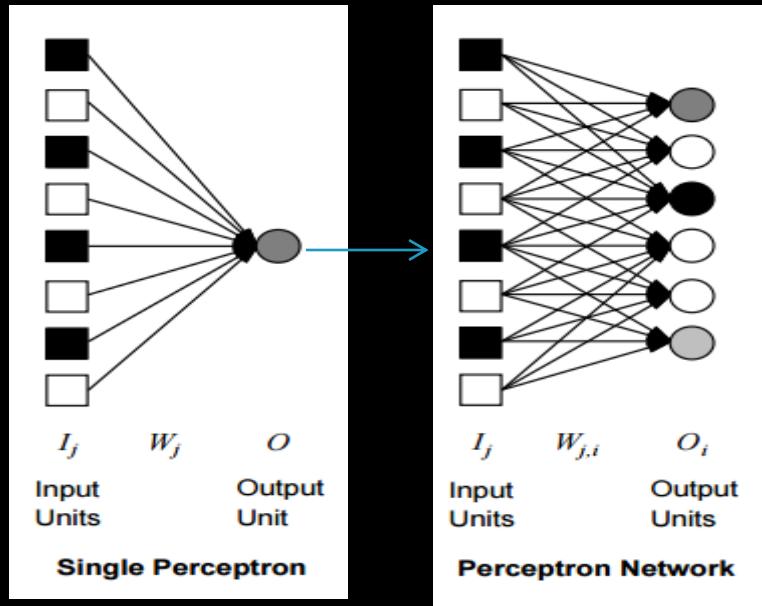


“有点难”

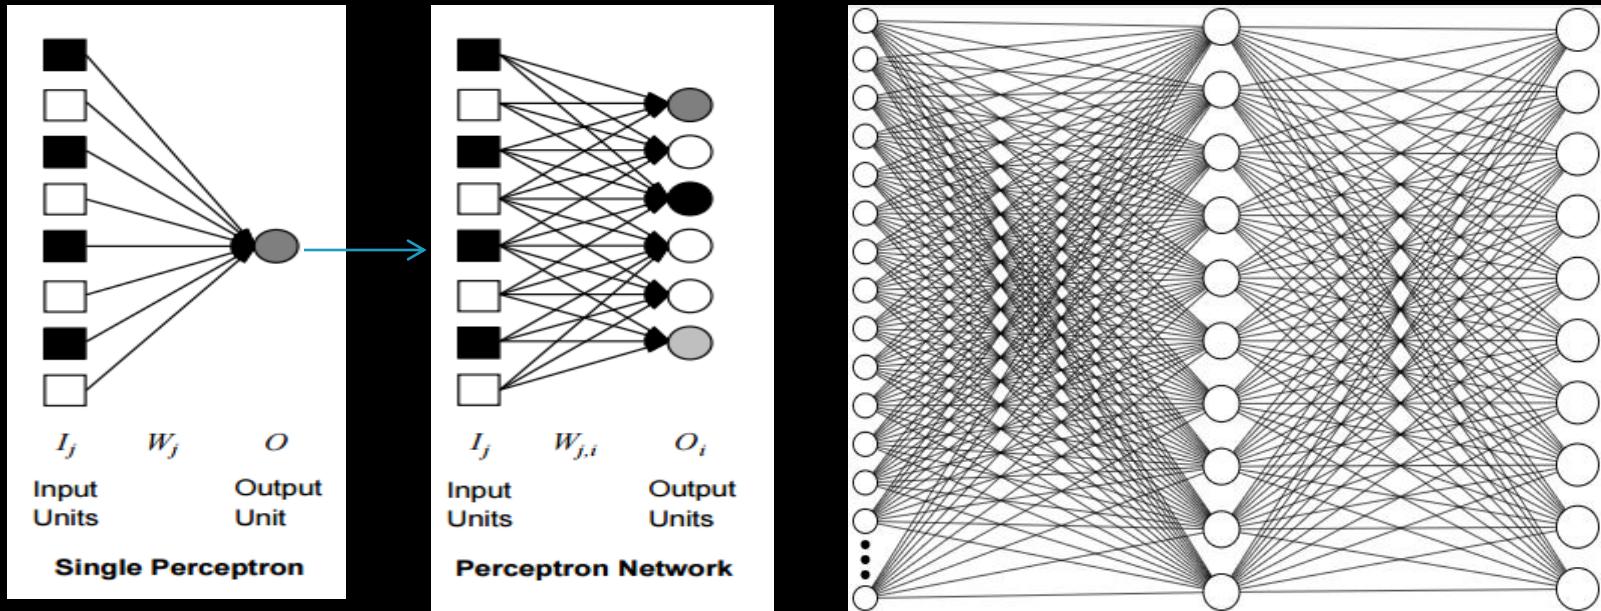
# 算法在手，层数你有



# 算法在手，层数你有



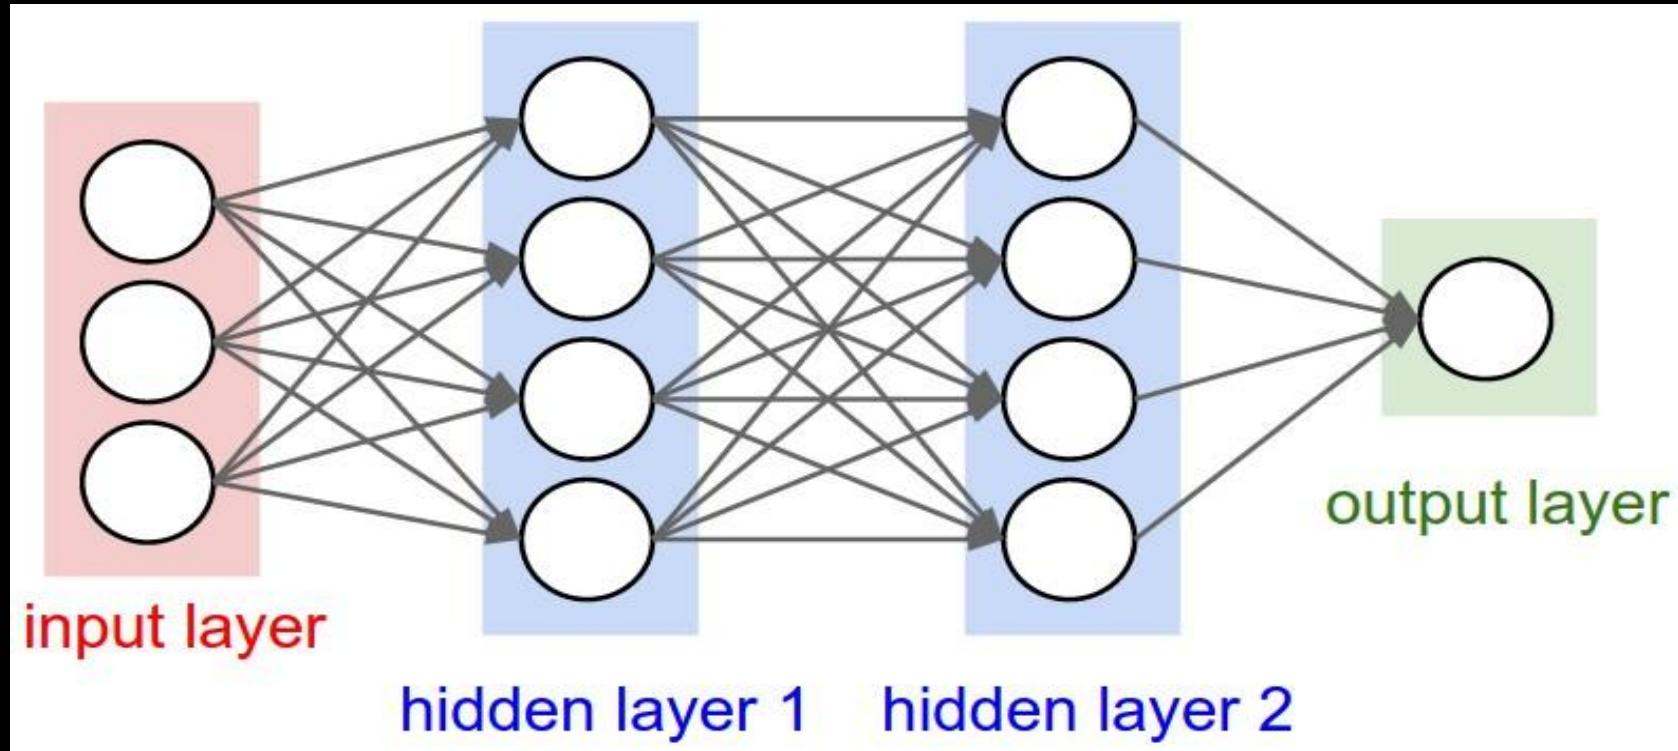
# 算法在手，层数你有



[https://en.wikipedia.org/wiki/Artificial\\_neural\\_network](https://en.wikipedia.org/wiki/Artificial_neural_network)

<https://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15381-f01/www/handouts/110601.pdf>

算法在手，层数你有



算法在手，层数你。。。额，不行了

# 不行了：反向传播算法的局限性

梯度消失

局部最优

过学习

# 关于过学习 (Overfit) 和欠学习 (Underfit)

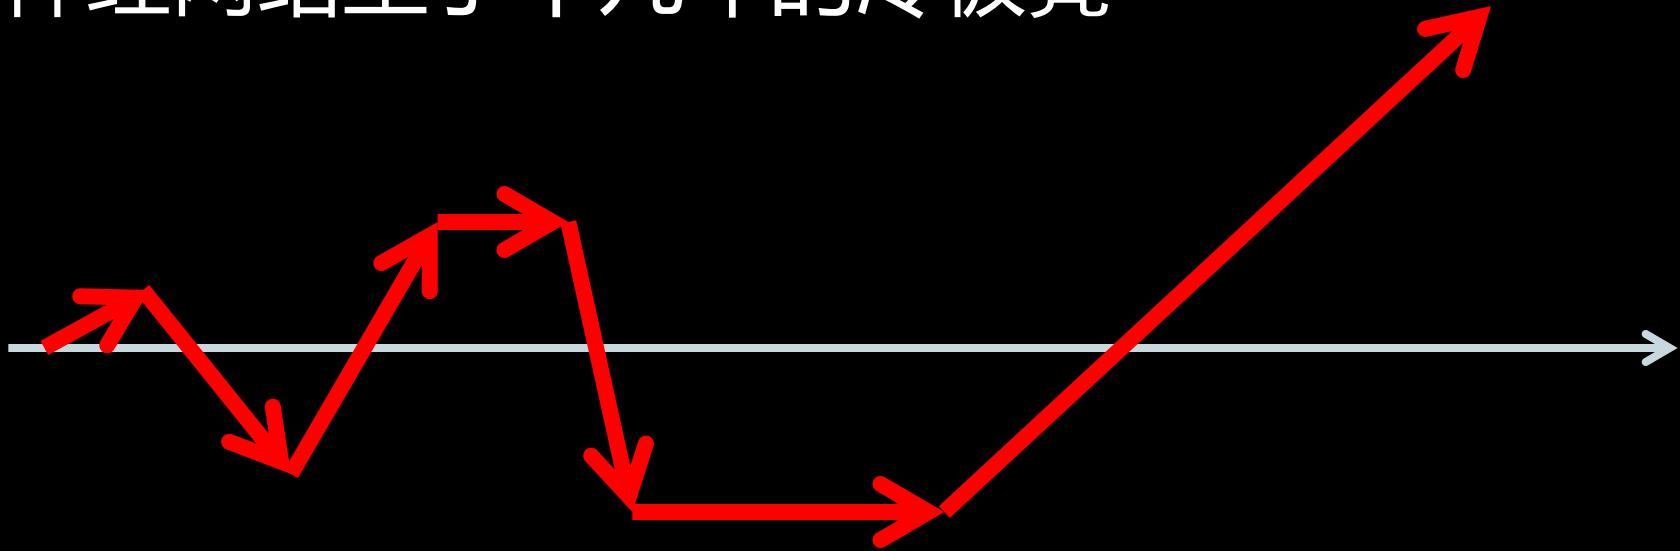
给我四个参数，我能拟合出来一头大象。

给我五个，我就能让大象的鼻子动起来。



—— 冯 • 诺依曼

神经网络坐了十几年的冷板凳

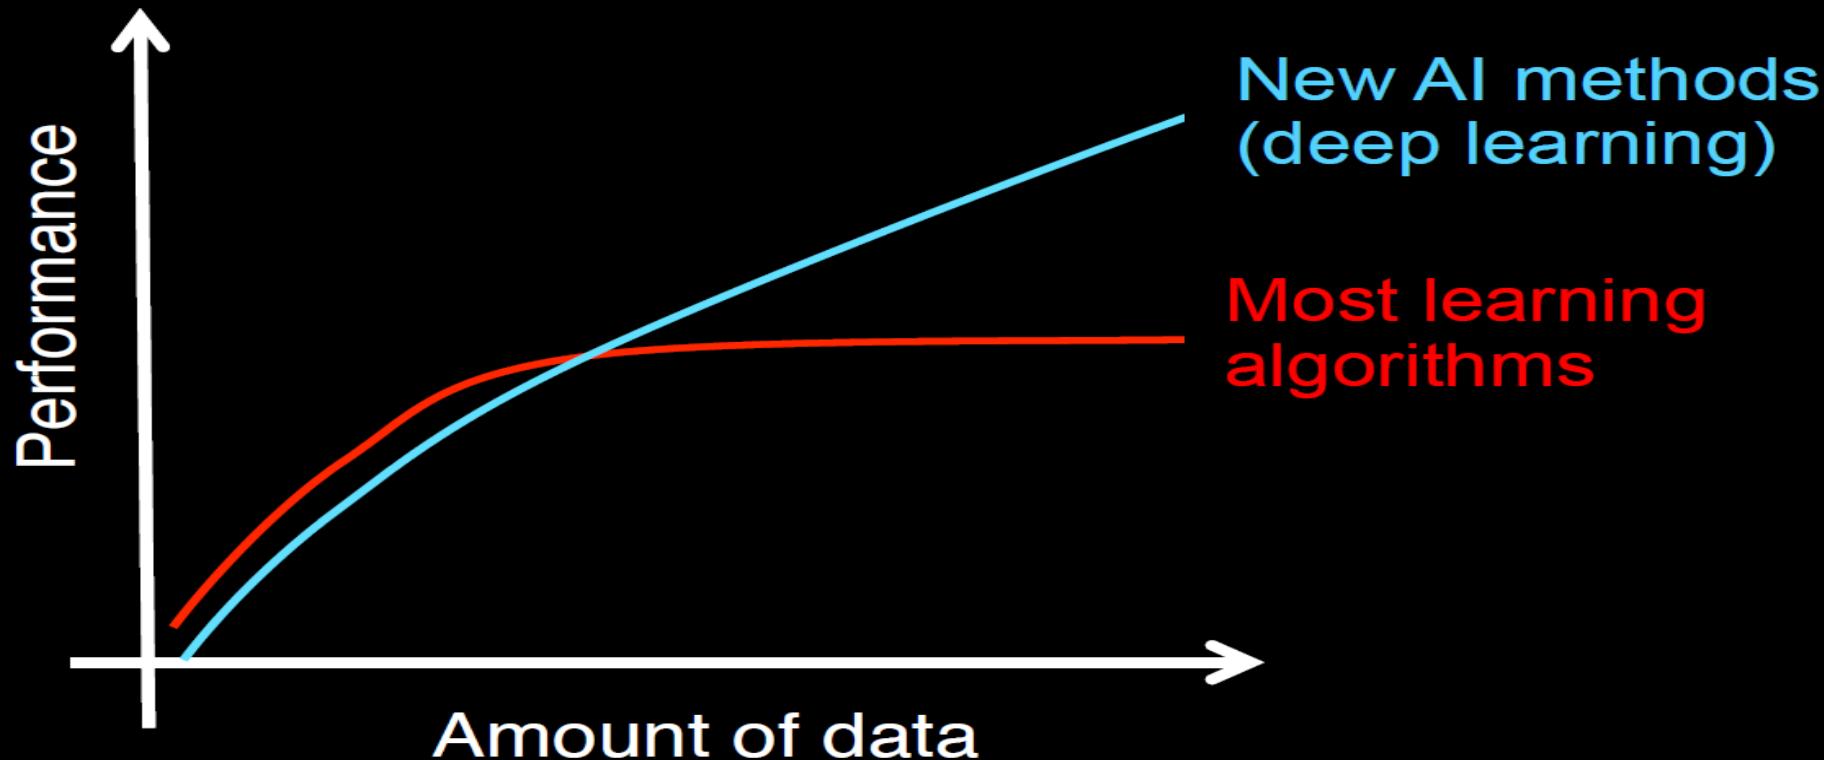


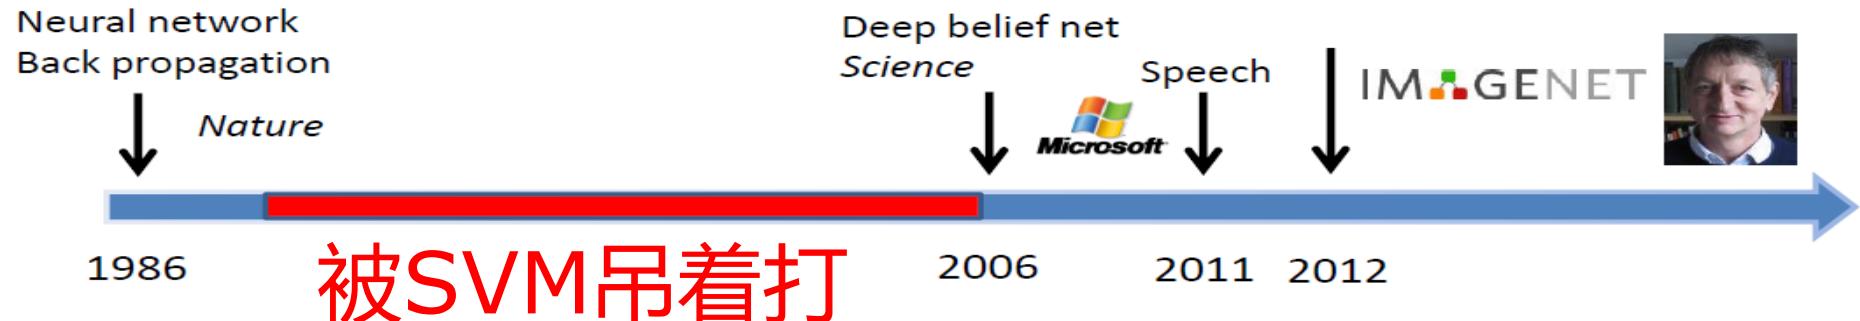
神经网络坐了十几年的冷板凳

被SVM吊打

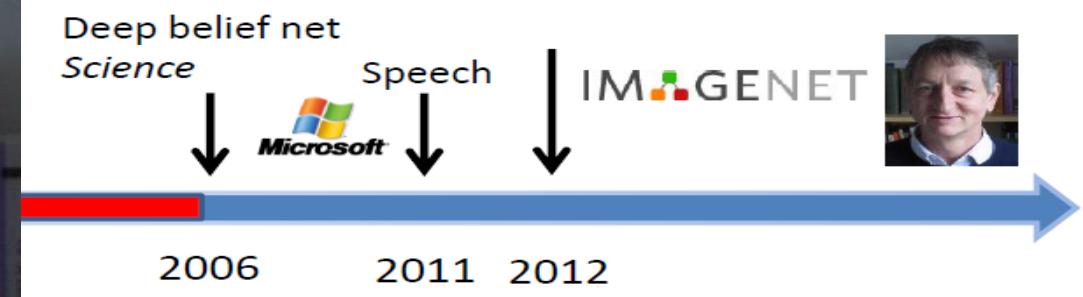


# Data and machine learning





Credit: Xiaoqiang Wang, The Chinese University of Hong Kong



面壁十年图破壁

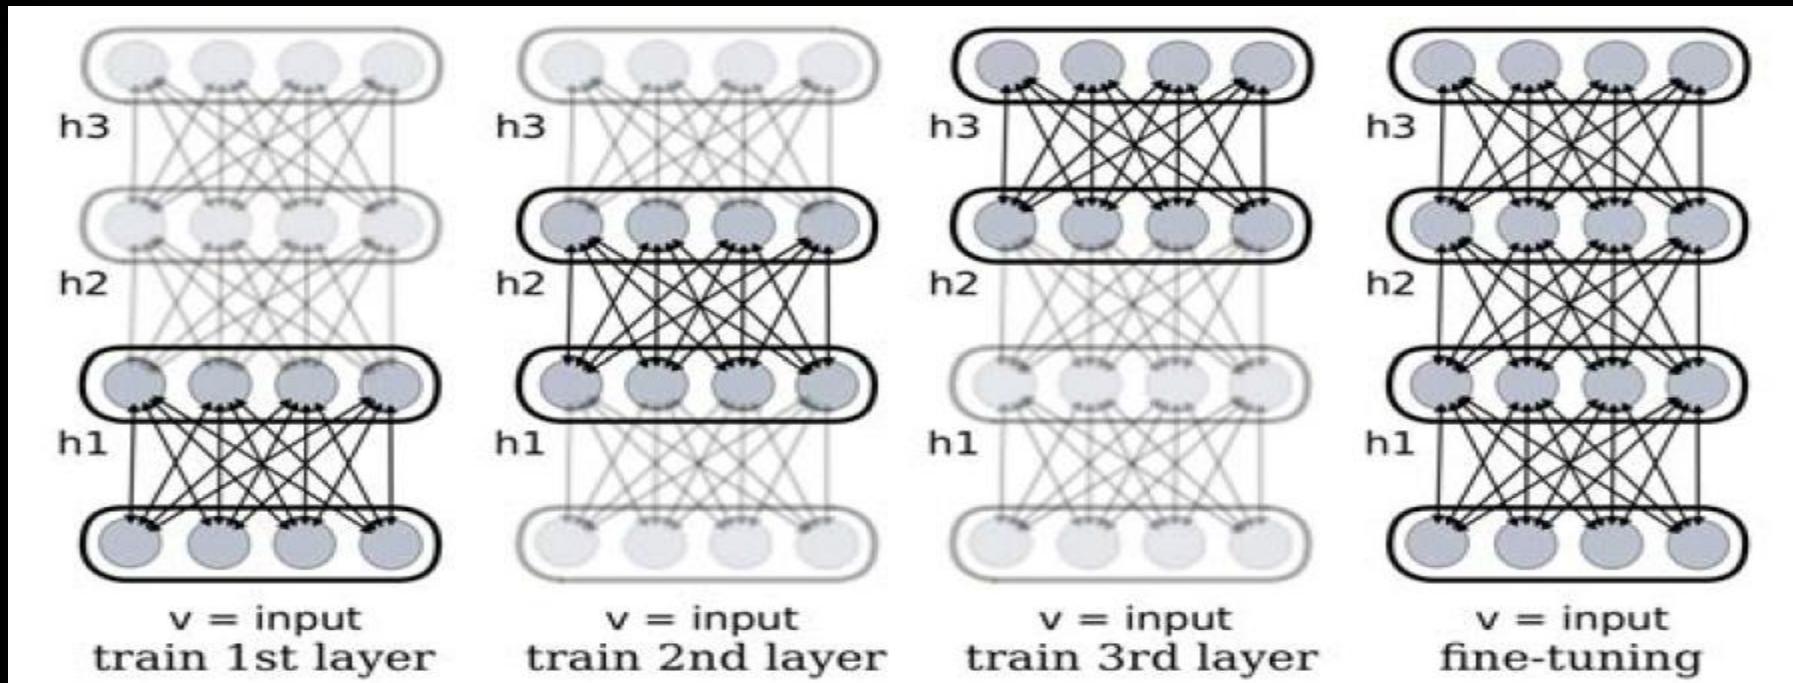
我Hinton悟到了

—！层！—！层！来！



Credit: Xiaoqiang Wang, The Chinese University of Hong Kong

# 一层一层先预训练参数，最后反向传播



Neural network  
Back propagation



1980



Deep belief net  
*Science*      *Speech*



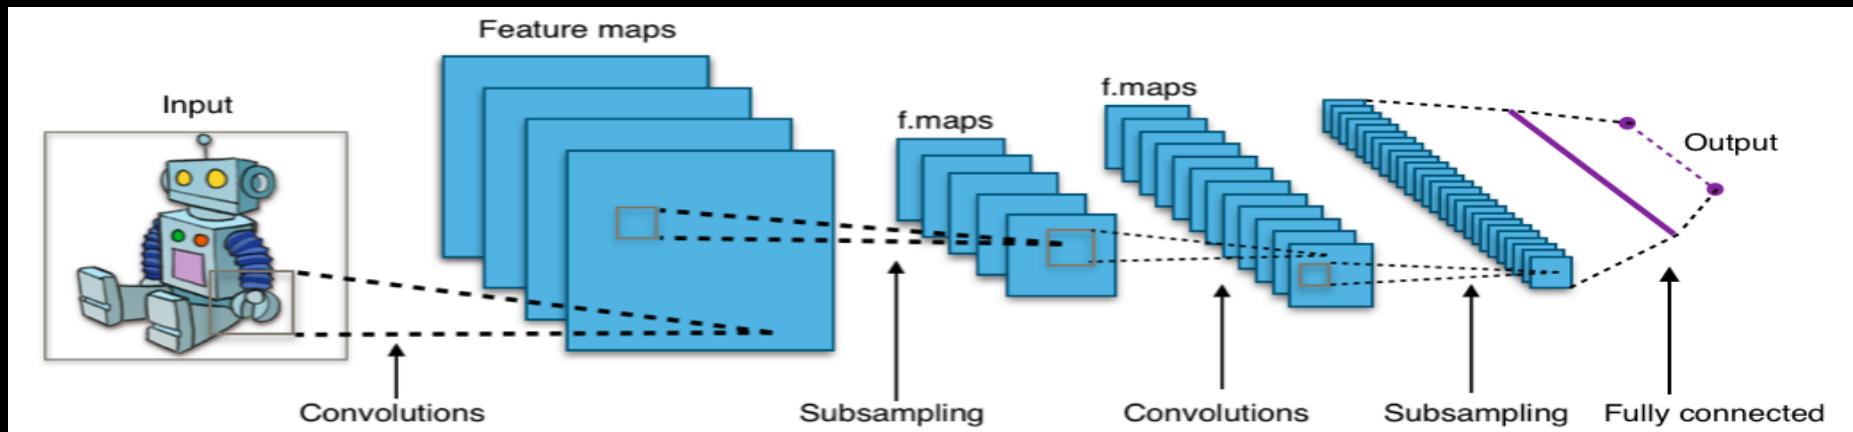
减少训练权重的规模

试试我乐康卷 ☺



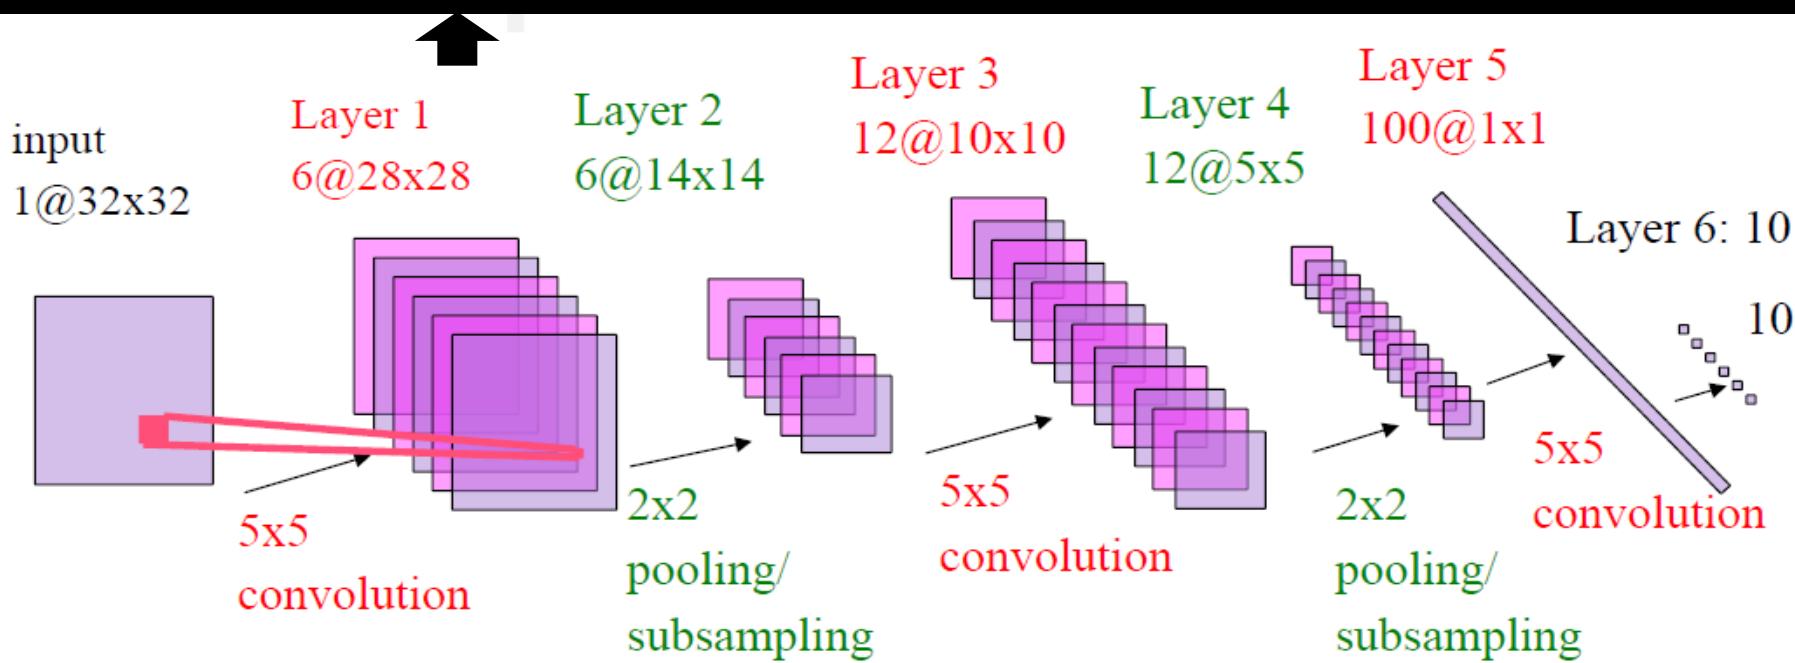
Credit: Xiaoqiang Wang, The Chinese University of Hong Kong

# 乐康卷：大体架构



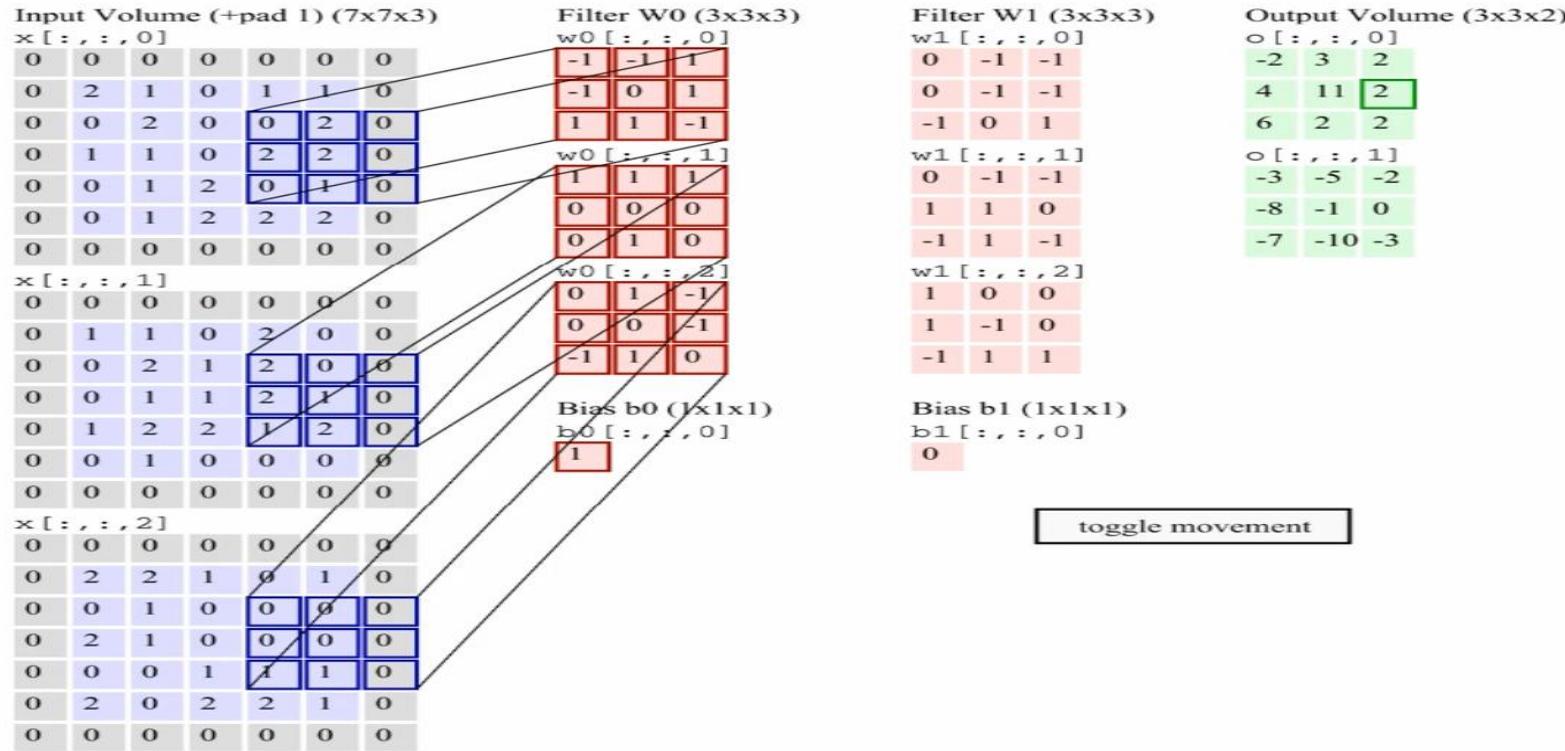
# 乐康卷：参数配置

## LeNet5



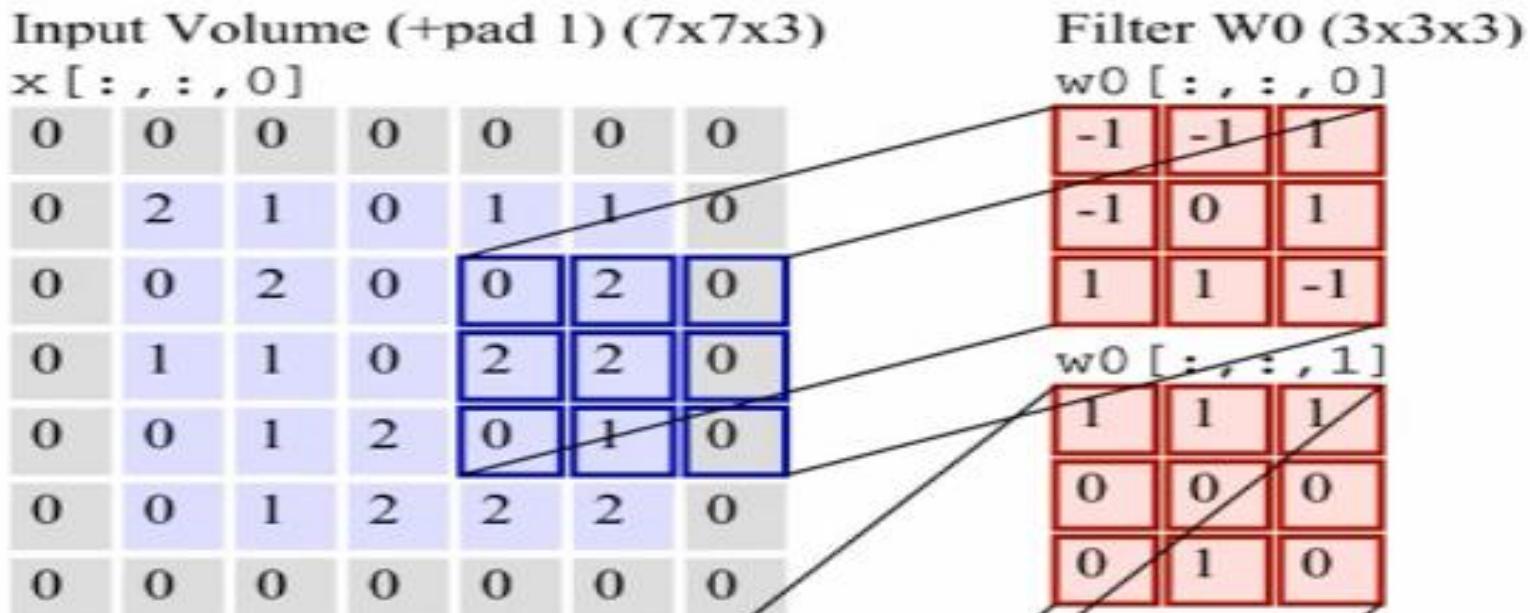
# 乐康卷：具体实例（来自LeCun大神）

Y LeCun



Animation: Andrej Karpathy <http://cs231n.github.io/convolutional-networks/>

# 乐康卷：具体实例（来自LeCun大神）



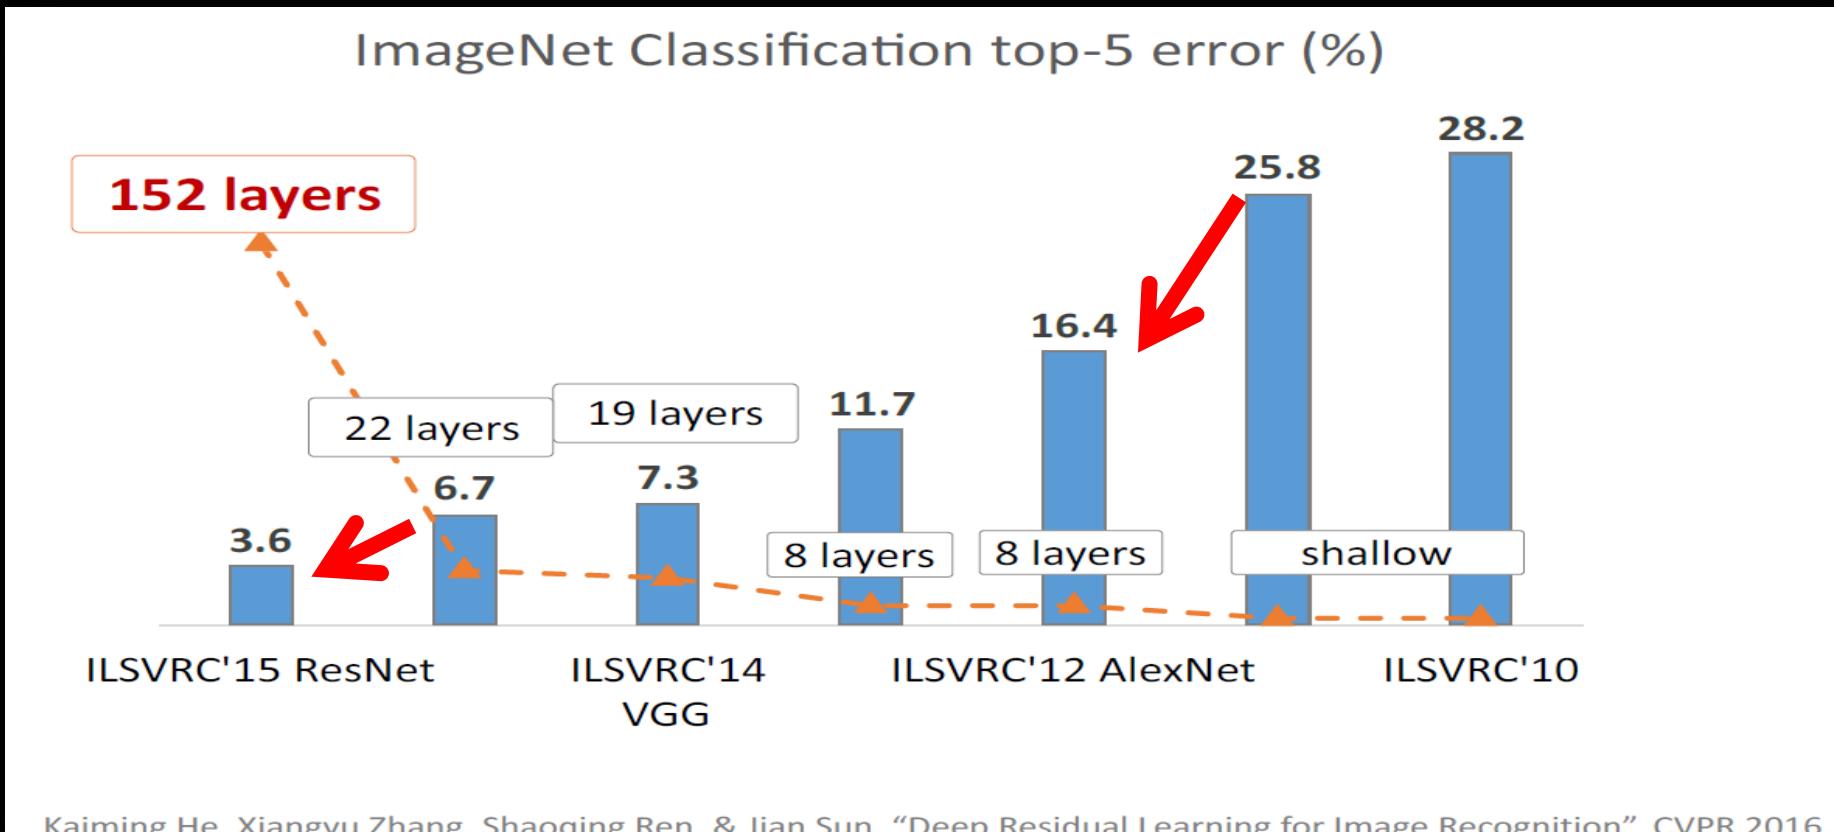
~~不行了：反向传播算法的局限性~~

梯度消失 - 逐层预训练

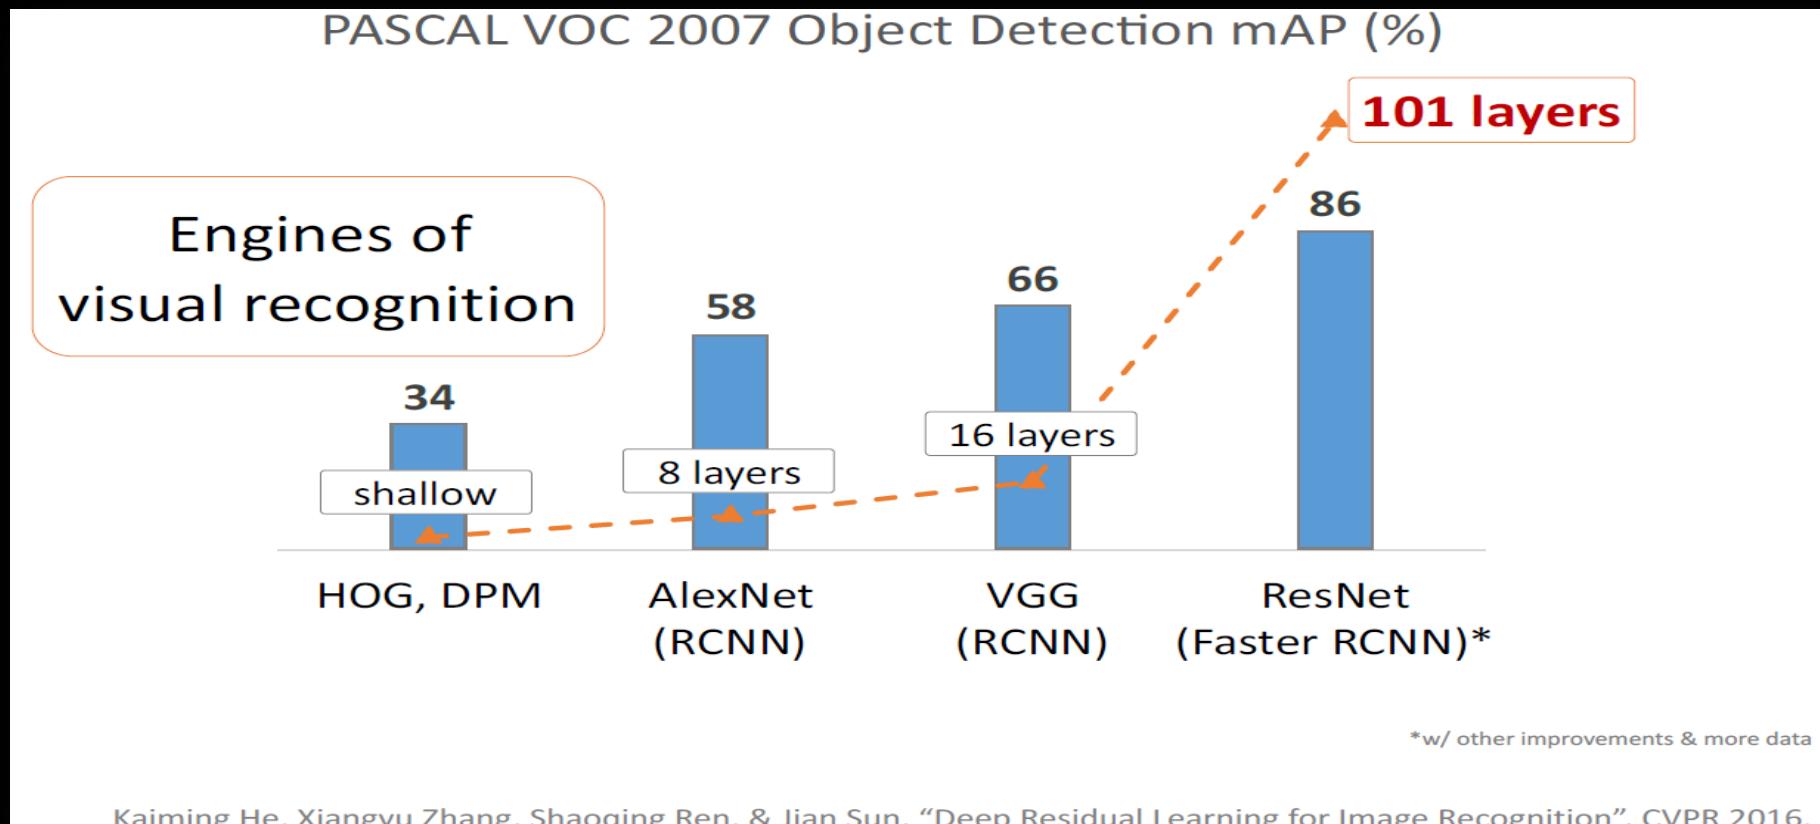
局部最优 - 多次随机化

过学习 - 大数据集

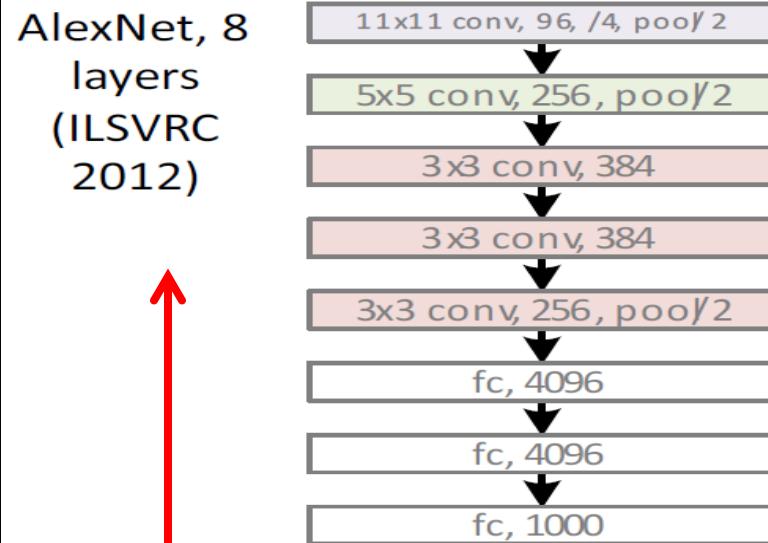
# 算法和模型都完备后，PhD们的心开始躁动起来



# 算法和模型都完备后，PhD们的心开始躁动起来



# 算法和模型都完备后，PhD们的心开始躁动起来

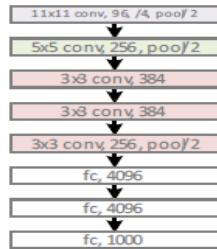


Caffe、mxnet中有现成模型

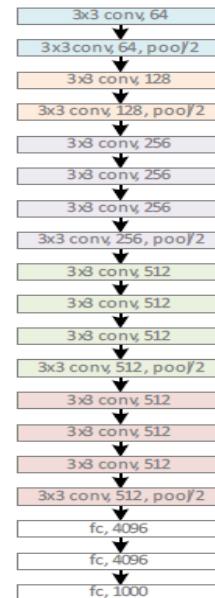
Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". CVPR 2016.

# 算法和模型都完备后，PhD们的心开始躁动起来

AlexNet, 8  
layers  
(ILSVRC  
2012)



VGG, 19  
layers  
(ILSVRC  
2014)



GoogleNet, 22  
layers  
(ILSVRC 2014)



Kaiming He, Xiangyu Zhang, Shaoqing Ren, & Jian Sun. "Deep Residual Learning for Image Recognition". CVPR 2016.

# 算法和模型都完备后，PhD们的心开始躁动起来

## Revolution of Depth

AlexNet, 8  
layers  
(ILSVRC  
2012)



VGG, 19  
layers  
(ILSVRC  
2014)



ResNet, 152  
layers  
(ILSVRC 2015)



# 算法和模型都完备后，PhD们的心开始躁动起来

README.md

## Deep Residual Networks with 1K Layers

By Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun.

Microsoft Research Asia (MSRA).

### Table of Contents

- 1. Introduction
- 2. Notes
- 3. Usage

### Introduction

This repository contains re-implemented code for the paper "Identity Mappings in Deep Residual Networks" (<http://arxiv.org/abs/1603.05027>). This work enables training quality 1k-layer neural networks in a super simple way.

*Acknowledgement:* This code is re-implemented by Xiang Ming from Xi'an Jiaotong University for the ease of release.

*See Also:* Re-implementations of ResNet-200 [a] on ImageNet from Facebook AI Research (FAIR):  
<https://github.com/facebook/fb.resnet.torch/tree/master/pretrained>

1001层DRN

# 故事背景

赵总看到这里很激动：之所以预测不好，  
一定是因为层数不够。直接加到两千层试试！

小李：.....

## 小结与思考

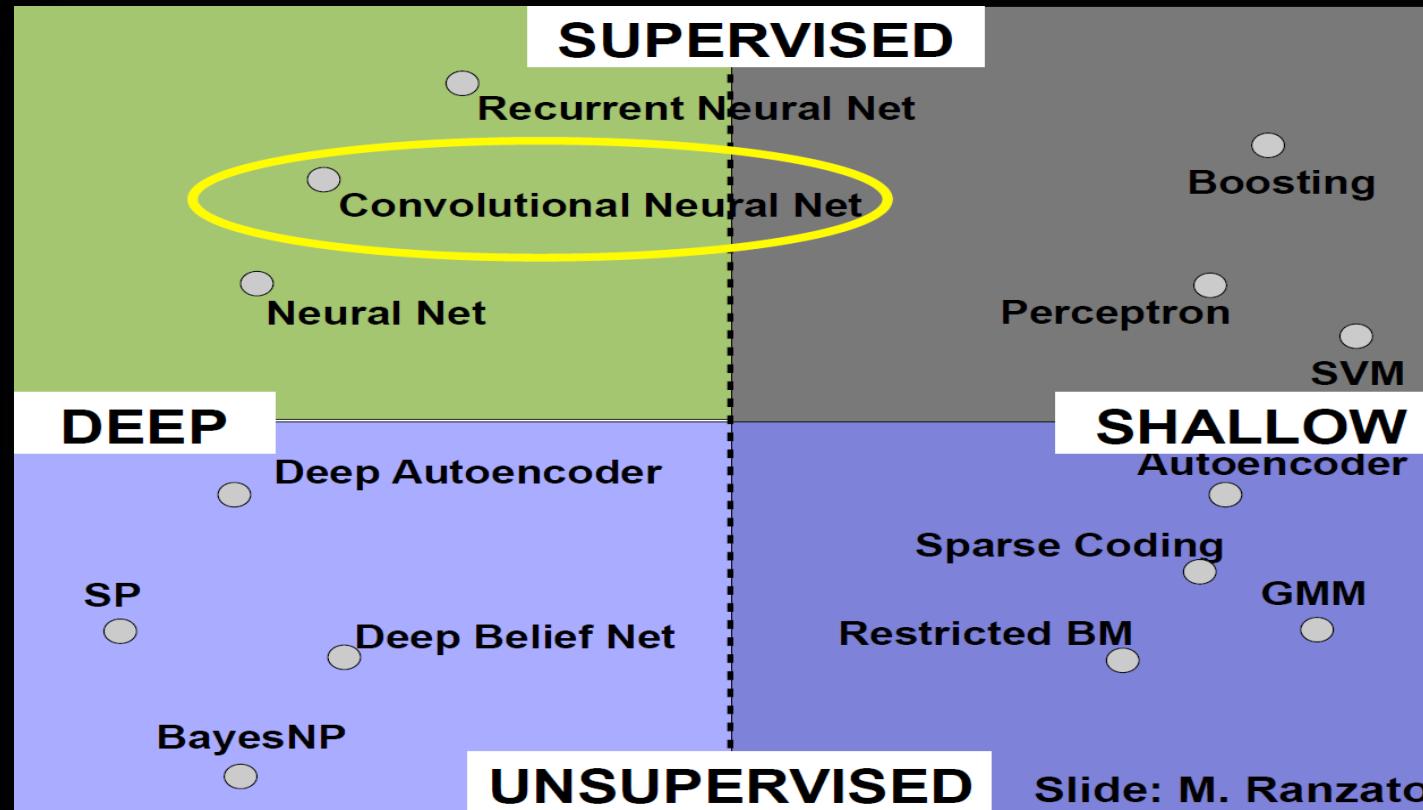
为什么深度学习会在图像、音频、NLP领域成功？

单纯的加层就可以提高预测的效果么？

扩大训练数据的数据规模就可以提高预测的效果么？

最重要的，如果一开始就**针对的错误的问题**.....

# 没有提到的跟机器学习有关的内容



# 没有提到的跟机器学习有关的内容

$$p(\mathbf{v}) = \sum_{\mathbf{h}} \frac{e^{-E(\mathbf{v}, \mathbf{h})}}{\sum_{\mathbf{u}, \mathbf{g}} e^{-E(\mathbf{u}, \mathbf{g})}}. \quad (1)$$

$$E(\mathbf{v}, \mathbf{h}) = - \sum_i a_i v_i - \sum_j b_j h_j - \sum_{i,j} w_{ij} v_i h_j, \quad (2)$$

$$p(\mathbf{v}|\mathbf{h}) = \prod_i p(v_i|\mathbf{h}) \quad \text{and} \quad p(v_i = 1|\mathbf{h}) = \text{sigm} \left( a_j + \sum_j h_j w_{ij} \right),$$

$$p(\mathbf{h}|\mathbf{v}) = \prod_j p(h_j|\mathbf{v}) \quad \text{and} \quad p(h_j = 1|\mathbf{v}) = \text{sigm} \left( b_j + \sum_i v_i w_{ij} \right), \quad (3)$$

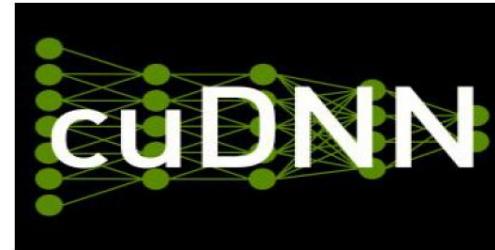
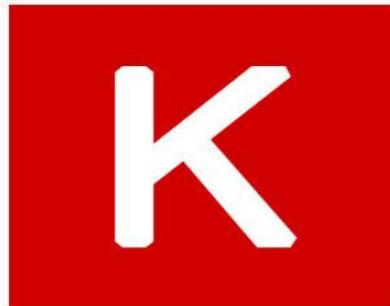
$$p(v_i = x|\mathbf{h}) = \frac{1}{\sigma_i \sqrt{2\pi}} \cdot e^{-\frac{(x - a_i - \sigma_i \sum_j w_{ij} h_j)^2}{2\sigma_i^2}},$$
$$p(h_j = 1|\mathbf{v}) = \text{sigm} \left( b_j + \sum_i \frac{v_i}{\sigma_i} w_{ij} \right). \quad (4)$$

想要学习深度学习，如何上手？

在2016年上手学深度学习是非常容易的

# Deep-Learning Package Zoo

- Torch
- Caffe
- Theano (Keras, Lasagne)
- CuDNN
- Tensorflow
- Mxnet
- Etc.



theano  
dmlc  
**mxnet**



入门阶段已经不需要自己写算法，写模型即可

## Deep-Learning Package Design Choices

- Model specification: **Configuration file** (e.g. Caffe, DistBelief, CNTK) versus **programmatic generation** (e.g. Torch, Theano, Tensorflow)
- For programmatic models, choice of high-level language: Lua (Torch) vs. Python (Theano, Tensorflow) vs others.
- We chose to work with **python** because of rich community and library infrastructure.

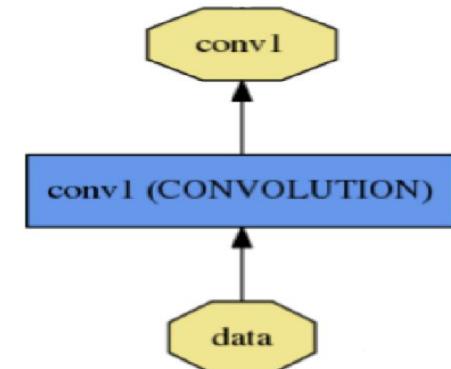
# 选取目前最常用的两个包进行介绍



# A Caffe Layer

```
name: "conv1"
type: CONVOLUTION
bottom: "data"
top: "conv1"
convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
        type: "xavier"
    }
}
```

} name, type, and the connection structure (input blobs and output blobs)  
layer-specific parameters

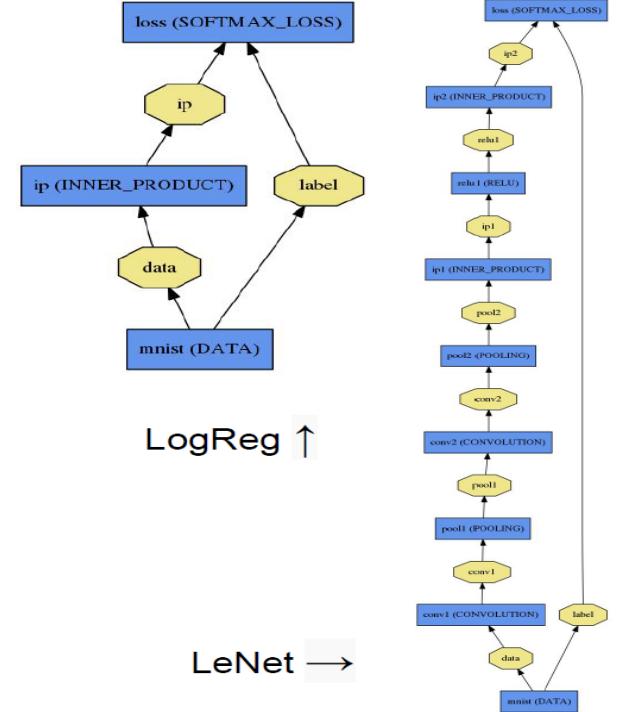


\*Link to the [Google Protobuffer Documentation](#)

# A Caffe Network

- A network is a set of layers connected as a DAG:

```
name: "dummy-net"  
  
layers { name: "data" ... }  
layers { name: "conv" ... }  
layers { name: "pool" ... }  
... more layers ...  
layers { name: "loss" ... }
```



# MNIST : 深度学习领域的 Hello World

## THE MNIST DATABASE

### of handwritten digits

[Yann LeCun](#), Courant Institute, NYU

[Corinna Cortes](#), Google Labs, New York

[Christopher J.C. Burges](#), Microsoft Research, Redmond

The MNIST database of handwritten digits, available from this page, has a training set of 60.000 examples, and a test set of 10.000 examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.

It is a good database for people who want to try learning techniques and pattern recognition methods on real-world data while spending minimal efforts on preprocessing and formatting.

Four files are available on this site:

[train-images-idx3-ubyte.gz](#): training set images (9912422 bytes)

[train-labels-idx1-ubyte.gz](#): training set labels (28881 bytes)

[t10k-images-idx3-ubyte.gz](#): test set images (1648877 bytes)

[t10k-labels-idx1-ubyte.gz](#): test set labels (4542 bytes)

0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 3  
3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9

# Misclassified examples

4	3	2	1	5	4	2	3	6	7
4	8	7	5	7	6	3	2	3	4
8	5	4	3	0	9	9	6	5	1
9	2	1	3	3	9	6	5	6	8
4	7	9	4	2	9	4	9	9	9
7	4	8	3	8	6	8	3	3	9
1	9	6	0	6	7	0	1	4	2
2	8	4	7	7	6	9	6	5	5
6	9	7	2	5	7	1	6	5	0
8	2	9	1	3	8	7	4	5	5

# Caffe

Deep learning framework  
by the [BVLC](#)

Created by  
[Yangqing Jia](#)  
Lead Developer  
[Evan Shelhamer](#)

 [View On GitHub](#)

## Training LeNet on MNIST with Caffe

We will assume that you have Caffe successfully compiled. If not, please refer to the [Installation page](#). In this tutorial, we will assume that your Caffe installation is located at `CAFFE_ROOT`.

### Prepare Datasets

You will first need to download and convert the data format from the MNIST website. To do this, simply run the following commands:

```
cd $CAFFE_ROOT
./data/mnist/get_mnist.sh
./examples/mnist/create_mnist.sh
```

If it complains that `wget` or `gunzip` are not installed, you need to install them respectively. After running the script there should be two datasets, `mnist_train_lmdb`, and `mnist_test_lmdb`.

### LeNet: the MNIST Classification Model

Before we actually run the training program, let's explain what will happen. We will use the [LeNet](#) network, which is known to work well on digit classification tasks. We will use a slightly different version from the original LeNet implementation, replacing the sigmoid activations with Rectified Linear Unit (ReLU) activations for the neurons.

# Caffe

Deep learning framework  
by the [BVLC](#)

Created by  
[Yangqing Jia](#)  
Lead Developer  
[Evan Shelhamer](#)

[View On GitHub](#)

# Caffe Model Zoo

Lots of researchers and engineers have made Caffe models for different tasks with all kinds of architectures and data. These models are learned and applied for problems ranging from simple regression, to large-scale visual classification, to Siamese networks for image similarity, to speech and robotics applications.

To help share these models, we introduce the model zoo framework:

- A standard format for packaging Caffe model info.
- Tools to upload/download model info to/from Github Gists, and to download trained `.caffemodel` binaries.
- A central wiki page for sharing model info Gists.

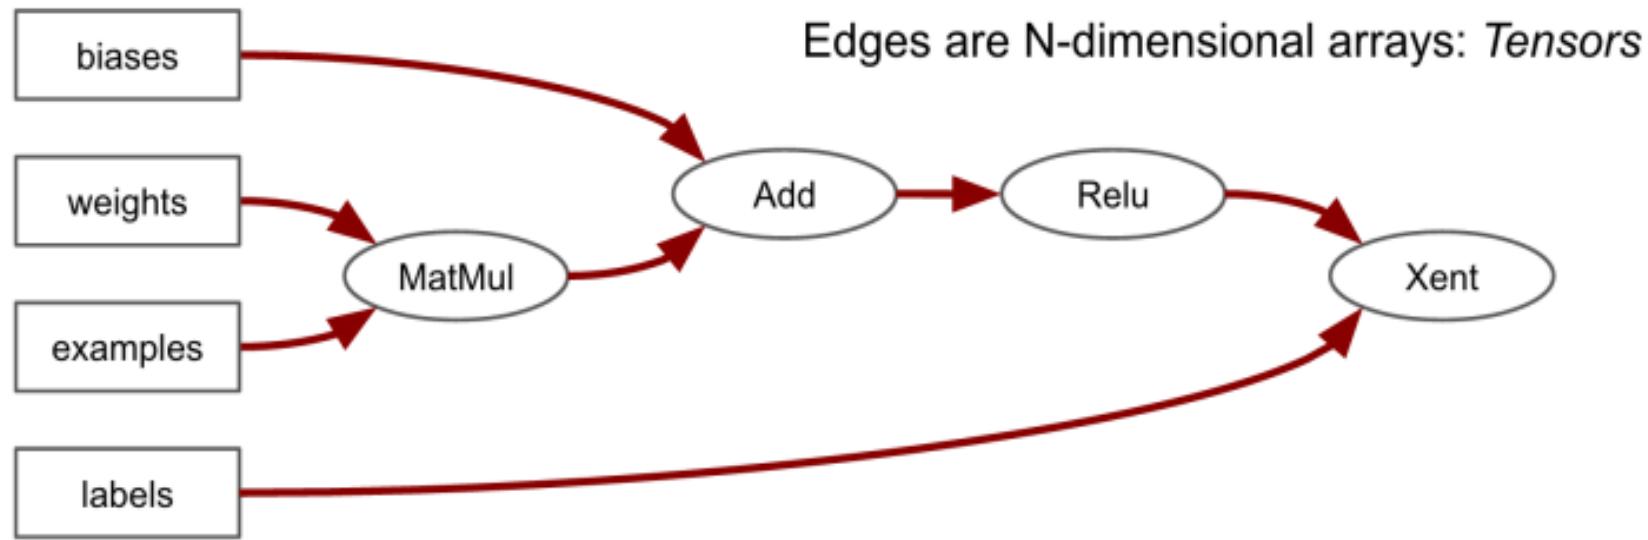
## Where to get trained models

First of all, we bundle BVLC-trained models for unrestricted, out of the box use.

See the [BVLC model license](#) for details. Each one of these can be downloaded by running `scripts/download_model_binary.py <dirname>` where `<dirname>` is specified below:

# TensorFlow : 搞清楚Tensor和Flow就行了

**TensorFlow is an open source software library for numerical computation using data flow graphs.**



# HELLO WORLD...

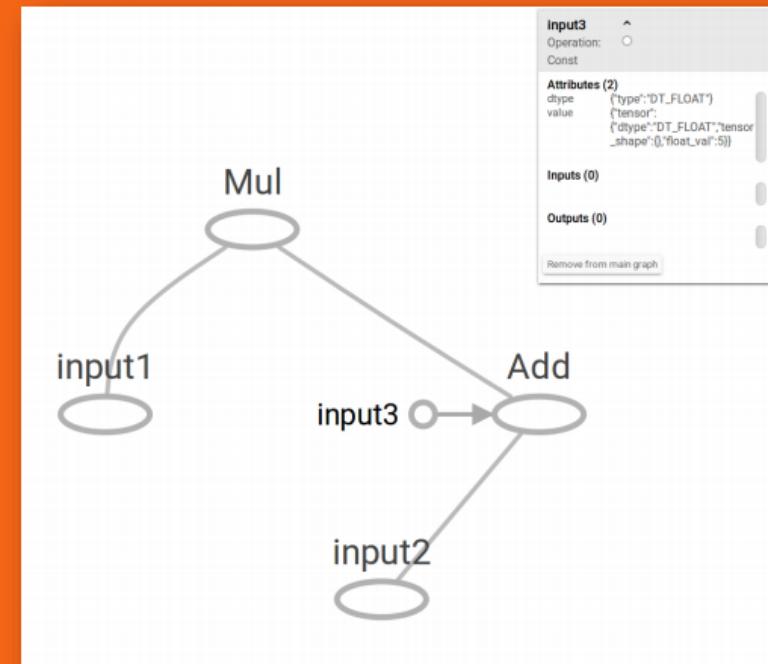
$$mul = i_1 * (i_2 + i_3)$$

```
import tensorflow as tf

input1 = tf.placeholder(tf.float32)
input2 = tf.placeholder(tf.float32)
input3 = tf.constant(5.0)
intermed = tf.add(input2, input3)
mul = tf.mul(input1, intermed)

with tf.Session() as sess:
    result = sess.run(
        [mul, intermed],
        feed_dict={input1:[1.0],
                   input2:[2.0]})

print(result)
```



Version: r0.11 ▾

## Basic Neural Networks

[MNIST For ML Beginners](#)

[About this tutorial](#)

[The MNIST Data](#)

[Softmax Regressions](#)

[Implementing the Regression](#)

[Training](#)

[Evaluating Our Model](#)

[Deep MNIST for Experts](#)

[About this tutorial](#)

[Setup](#)

[Load MNIST Data](#)

[Start TensorFlow InteractiveSession](#)

[Build a Softmax Regression Model](#)

[Placeholders](#)

[Variables](#)

[Predicted Class and Loss Function](#)

# MNIST For ML Beginners

This tutorial is intended for readers who are new to both machine learning and TensorFlow. If you already know what MNIST is, and what softmax (multinomial logistic) regression is, you might prefer this [faster paced tutorial](#). Be sure to [install TensorFlow](#) before starting either tutorial.

When one learns how to program, there's a tradition that the first thing you do is print "Hello World." Just like programming has Hello World, machine learning has MNIST.

MNIST is a simple computer vision dataset. It consists of images of handwritten digits like these:



It also includes labels for each image, telling us which digit it is. For example, the labels for the above images are 5, 0, 4, and 1.

In this tutorial, we're going to train a model to look at images and predict what digits they are. Our goal isn't to train a really elaborate model that achieves state-of-the-art performance -- although we'll give you code to do that later! -- but rather to dip a toe into using TensorFlow. As such, we're going to start with a very simple model, called a Softmax Regression.

The actual code for this tutorial is very short, and all the interesting stuff happens in just three lines. However, it is very important to understand the ideas behind it: both how TensorFlow works and the core machine learning concepts. Because of this, we are going to very carefully work through the code.

README.md

# TensorFlow Models

This repository contains machine learning models implemented in [TensorFlow](#). The models are maintained by their respective authors.

To propose a model for inclusion please submit a pull request.

## Models

- [autoencoder](#) -- various autoencoders
- [inception](#) -- deep convolutional networks for computer vision
- [namigner](#) -- recognize and generate names
- [neural\\_gpu](#) -- highly parallel neural computer
- [privacy](#) -- privacy-preserving student models from multiple teachers
- [resnet](#) -- deep and wide residual networks
- [slim](#) -- image classification models in TF-Slim
- [swivel](#) -- the Swivel algorithm for generating word embeddings
- [syntaxnet](#) -- neural models of natural language syntax
- [textsum](#) -- sequence-to-sequence with attention model for text summarization.
- [transformer](#) -- spatial transformer network, which allows the spatial manipulation of data within the network
- [im2txt](#) -- image-to-text neural network for image captioning.

深度学习怎么玩，取决于你多有钱

# GPU Acceleration

-gpu N flag tells caffe which gpu to use

Alternatively, specify solver\_mode: GPU in solver.prototxt

16.2s

10x

163s

Tesla K40m, cuDNN v3-RC  
ECC off, autoboot on

Intel® Xeon® E5-2698 v3 @  
2.60GHz, 3.6GHz turbo, (16  
cores total), HT off

Benchmark: Train Caffenet model, 20 iterations, 256x256 images, mini-batch size 256

基础 | 前沿 | 应用 | 评价

前沿 | NPU | GAN | 自编程

语言处理 | 语音识别 | 图像识别 | 视频理解

# 自然语言处理 | TL;DR

词性标注 | 组块识别 | 命名实体识别 | 语法解析 | 语义角色标注

乔姆斯基 | n-gram

# 自然语言处理 | RNN

# 自然语言处理 | LSTM, Bi-LSTM

# 自然语言处理 | Memory Network

自然语言处理 | DNC(?)

# 自然语言处理 | 接近并超越传统方法

<b>Task</b>		<b>Benchmark</b>	<b>SENNNA</b>
Part of Speech (POS)	(Accuracy)	97.24 %	97.29 %
Chunking (CHUNK)	(F1)	94.29 %	94.32 %
Named Entity Recognition (NER)	(F1)	89.31 %	89.59 %
Parse Tree level 0 (PT0)	(F1)	91.94 %	92.25 %
Semantic Role Labeling (SRL)	(F1)	77.92 %	75.49 %

机器自然语言翻译 | 大幅度改进

语音识别 | 热火朝天

语音识别 | 百度,微软,讯飞,谷歌,苹果

# 语音识别 | GMM-HMM (李开复)

# 语音识别 | Warp-CTC (百度)

Connectionist Temporal Classification

语音识别 | FSMN, DFCNN (讯飞)

前馈型序列记忆网络 | 深度全序列卷积神经网络

语音识别 | ResNet+ $\Sigma$  (微软)

NIST 2000 Switchboard WER 5.9%

TTS | WaveNet (谷歌)

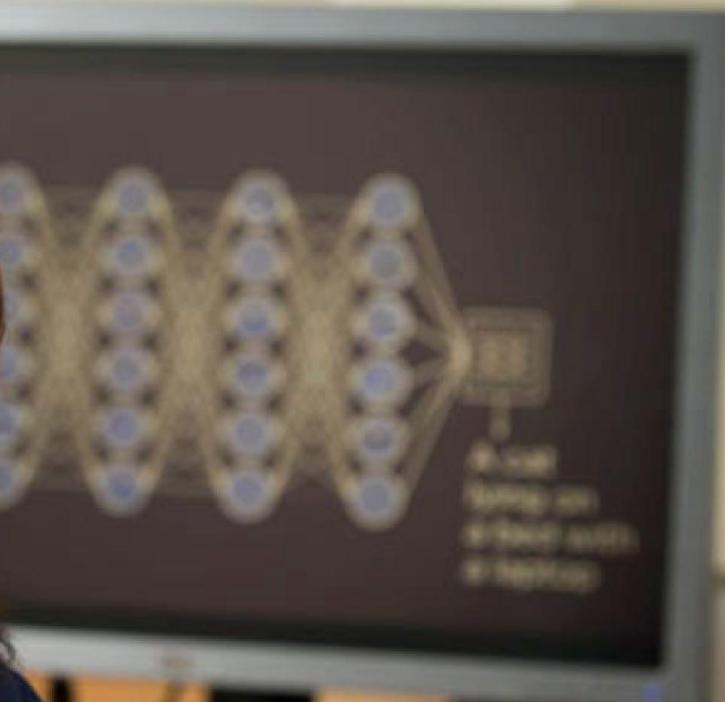
图像识别 | 成果最突出显著

# 图像识别 | MNIST

0 0 0 0 0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2 2 2 2 3  
3 3 3 3 3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9 9 9 9 9

图像识别 | IMGENET

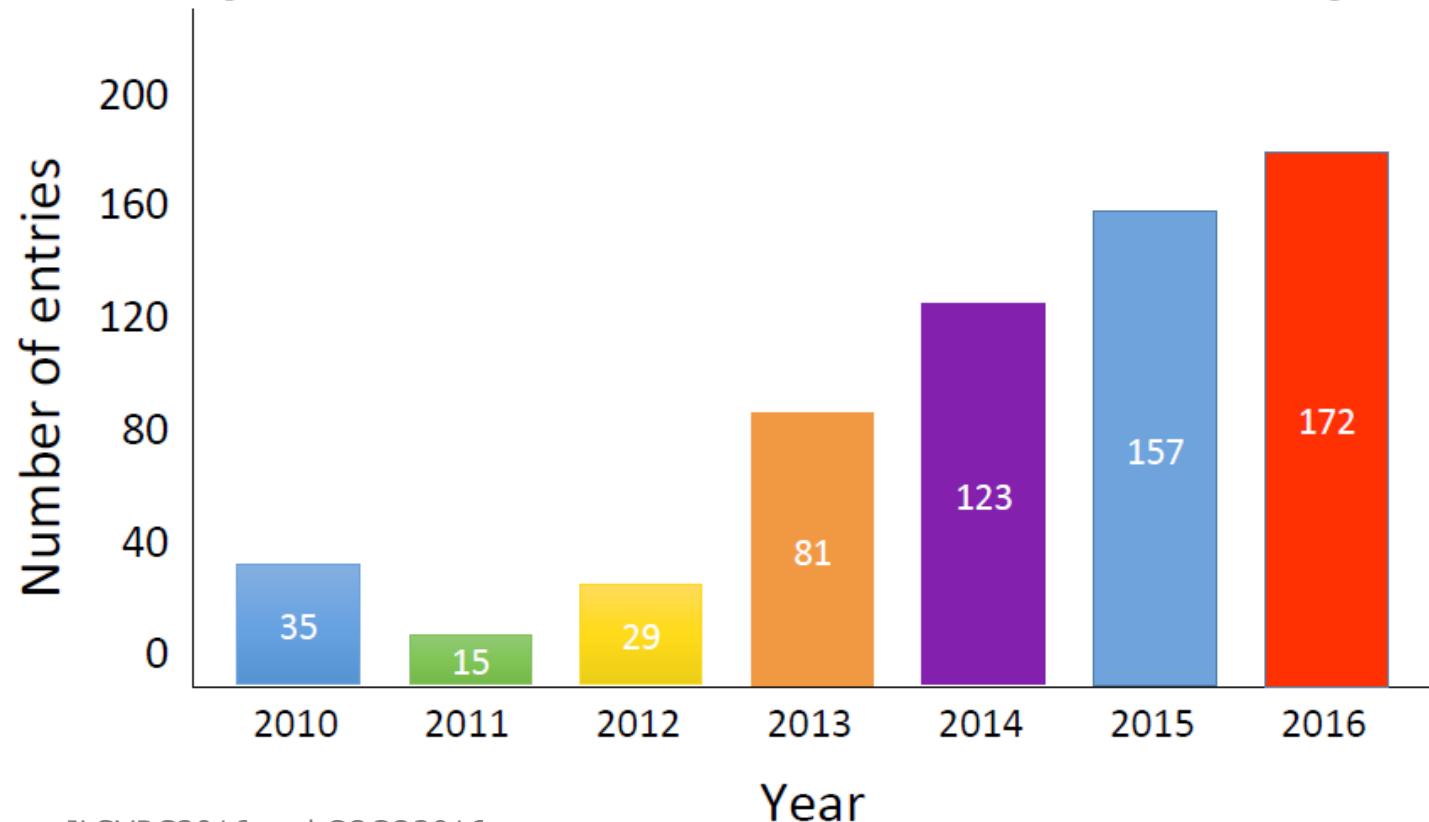
WordNet | 21K | 14M | 1M



# 图像识别 | ILSVRC

始于2010, 路碑

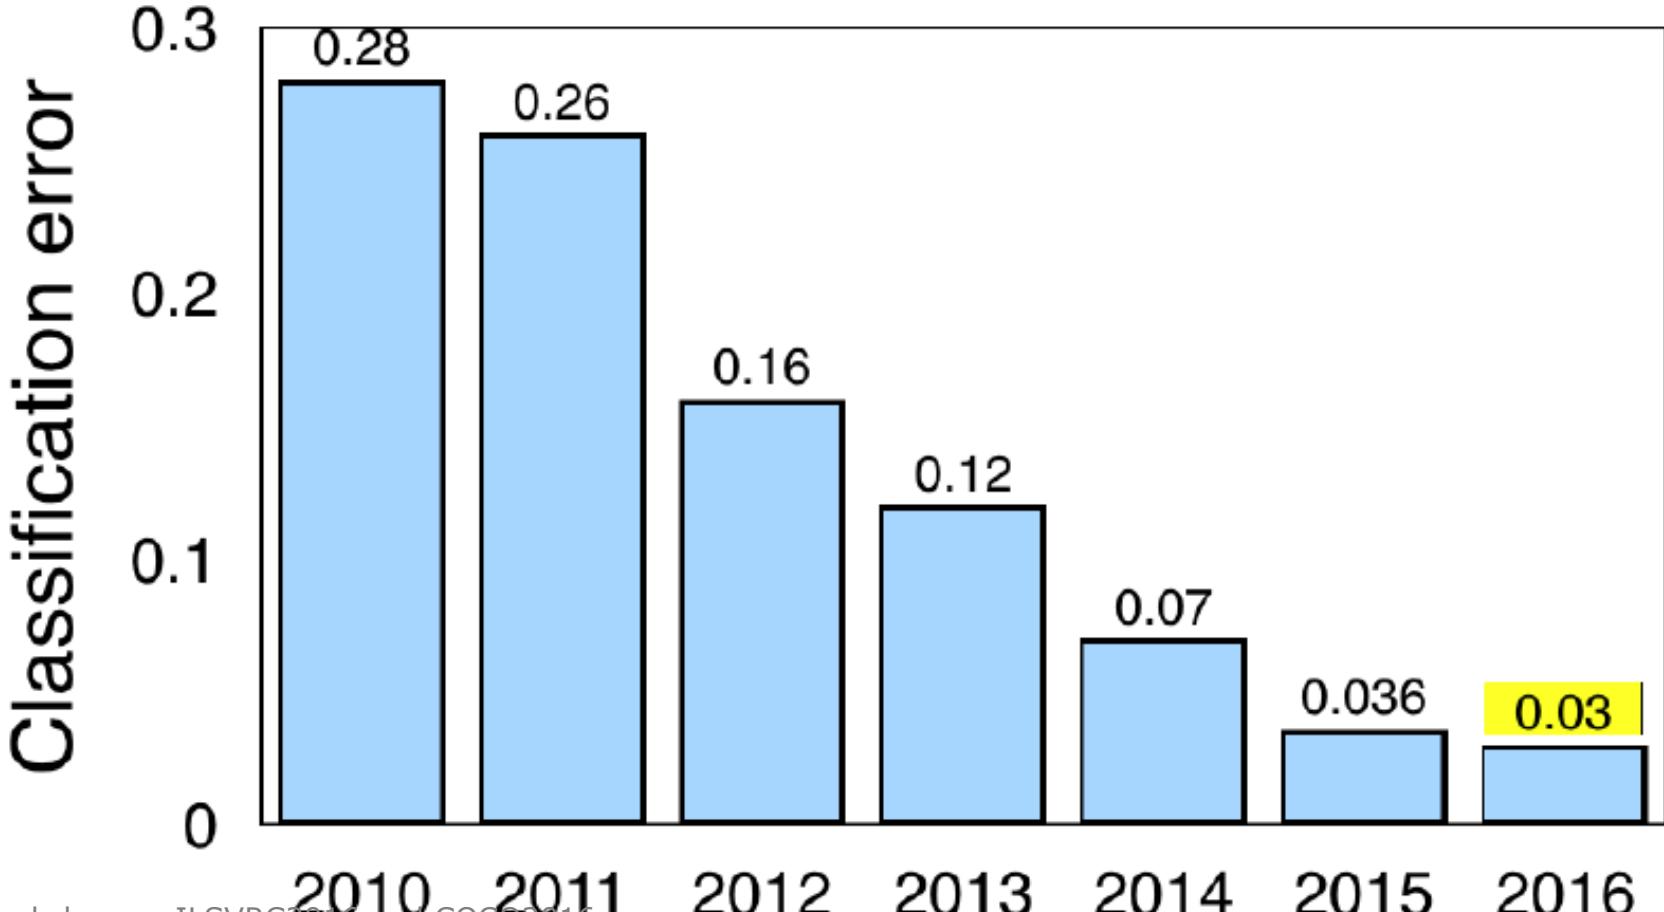
# Participation in ILSVRC over the years





credit: workshop on ILSVRC2016 and COCO2016

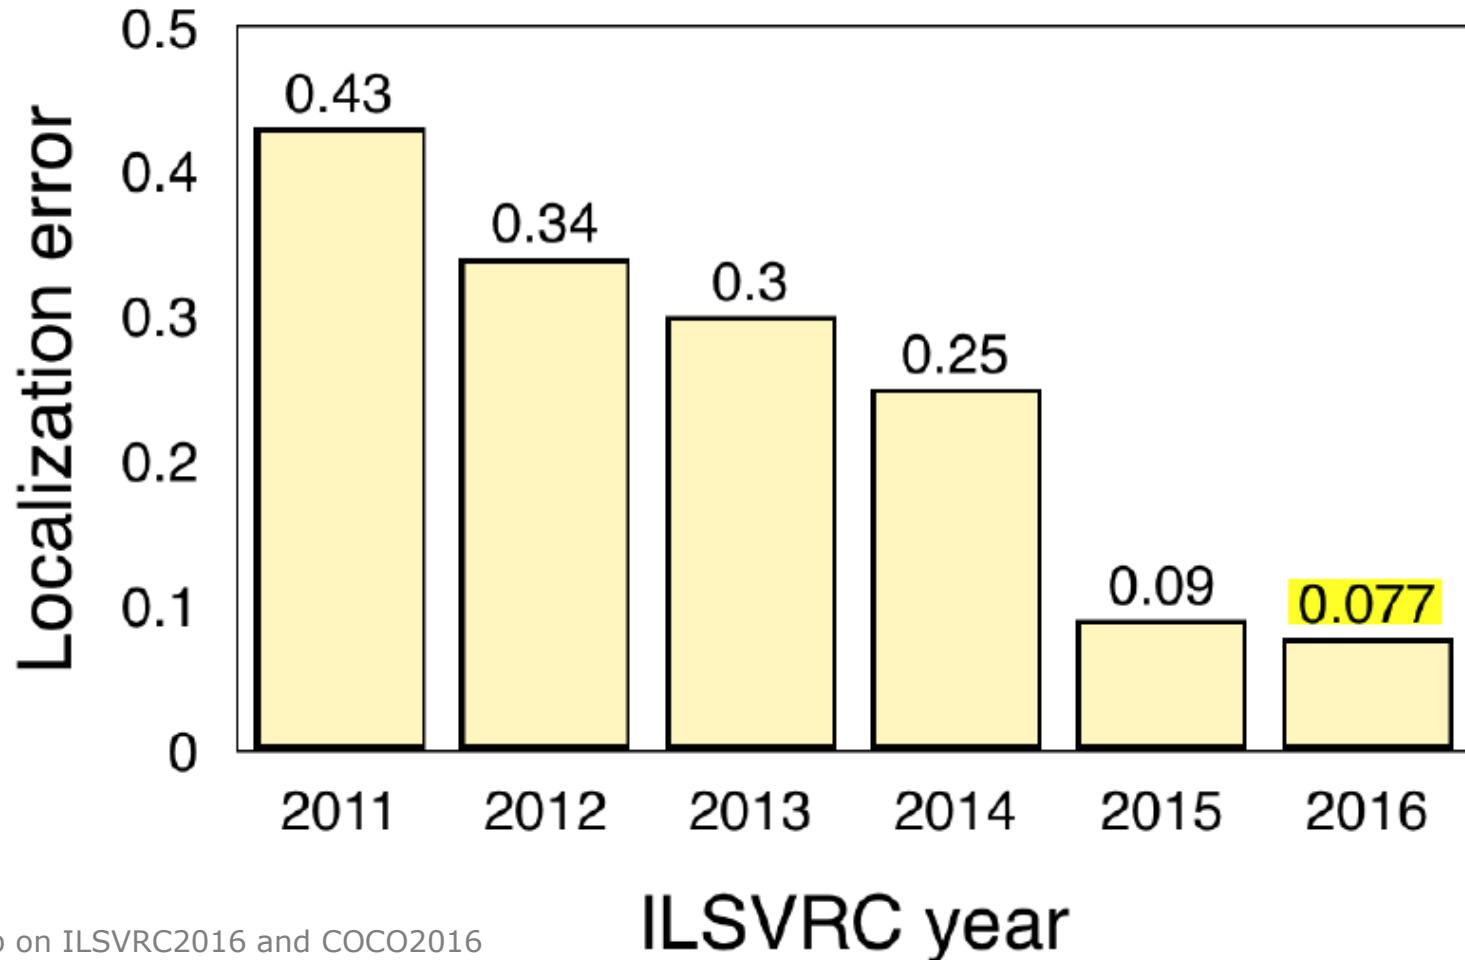
# Classification

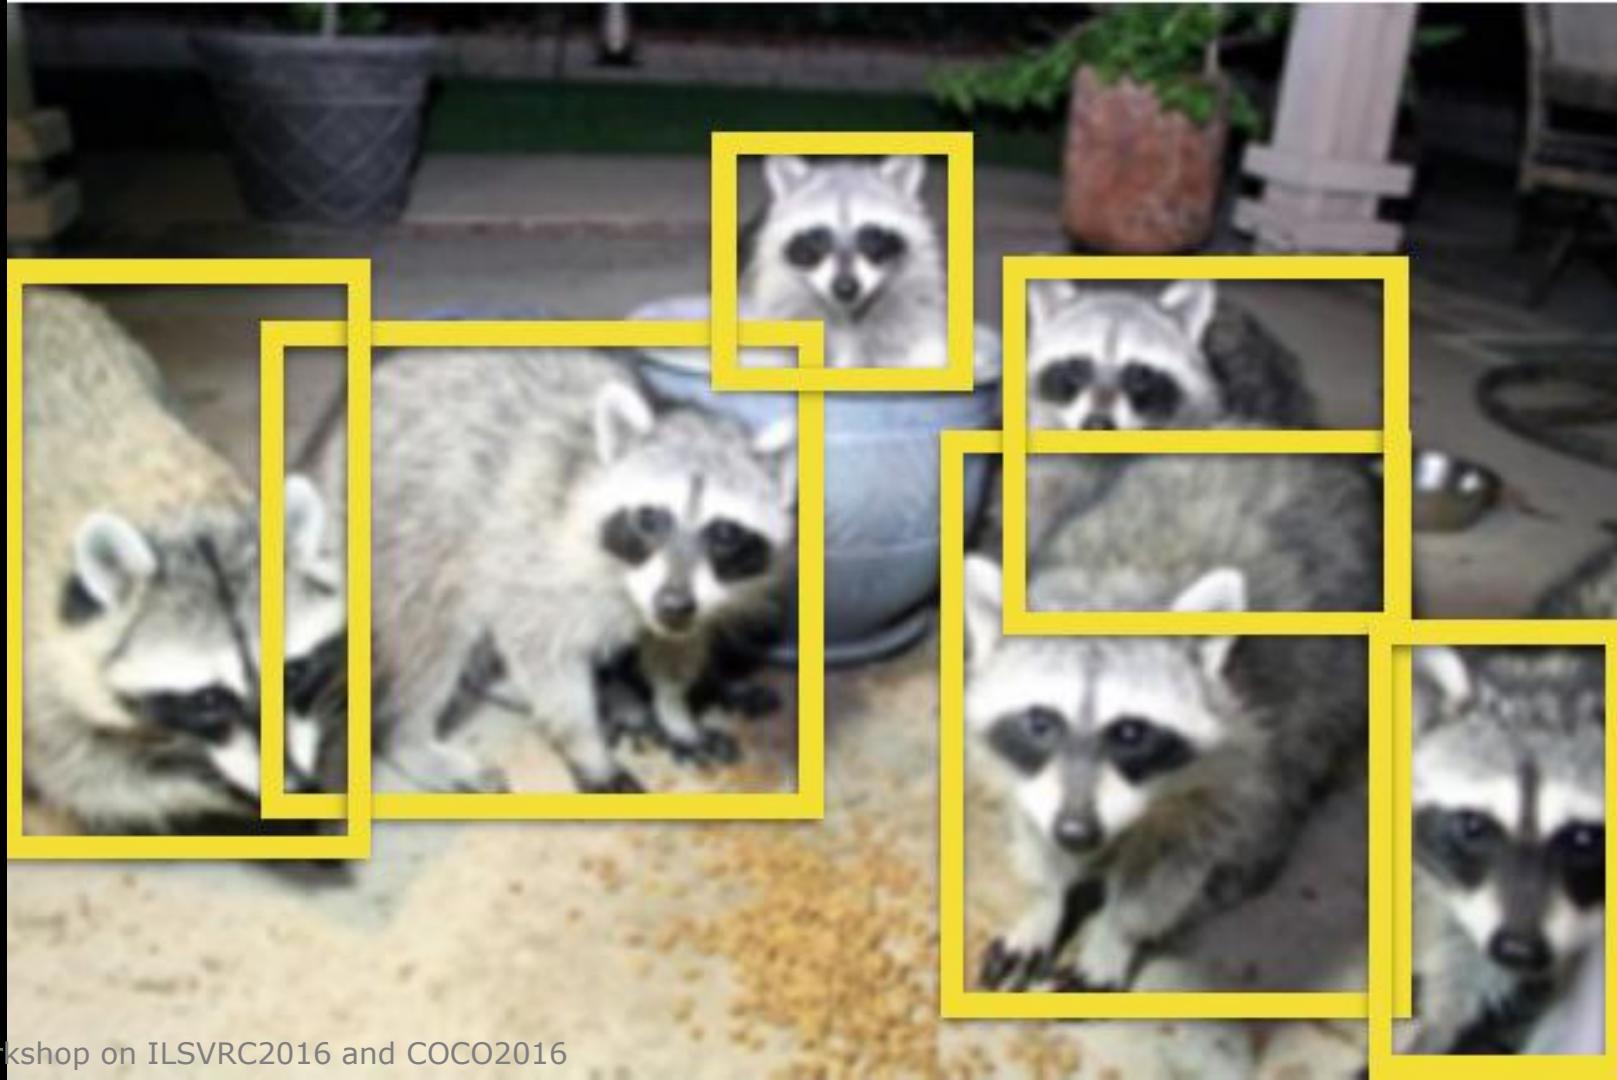




credit: workshop on ILSVRC2016 and COCO2016

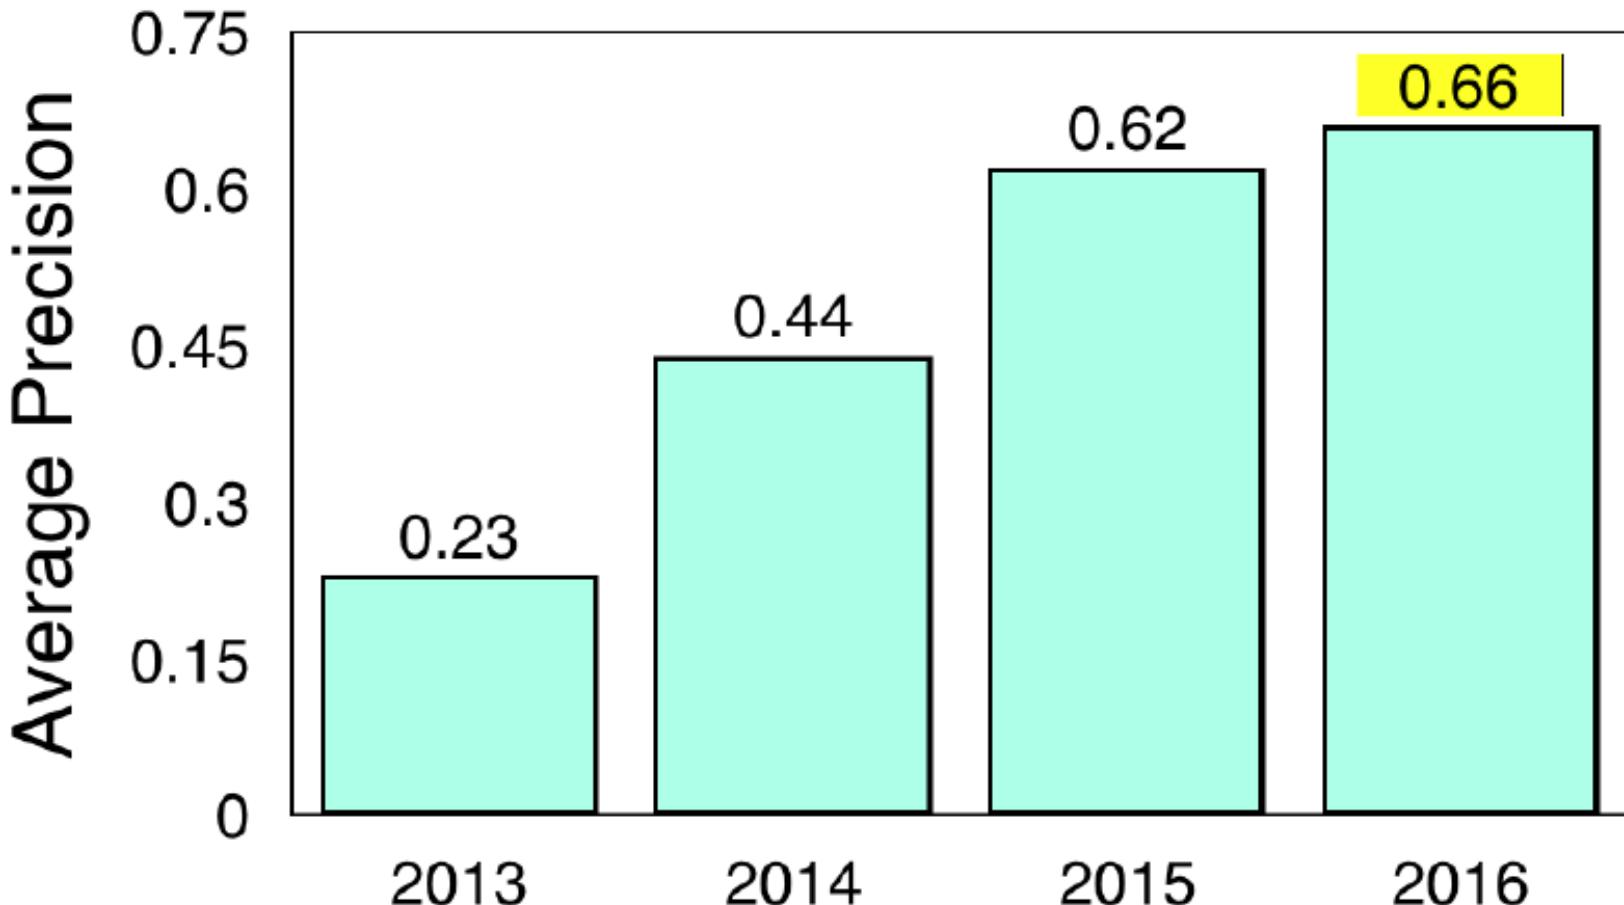
# Localization





credit: workshop on ILSVRC2016 and COCO2016

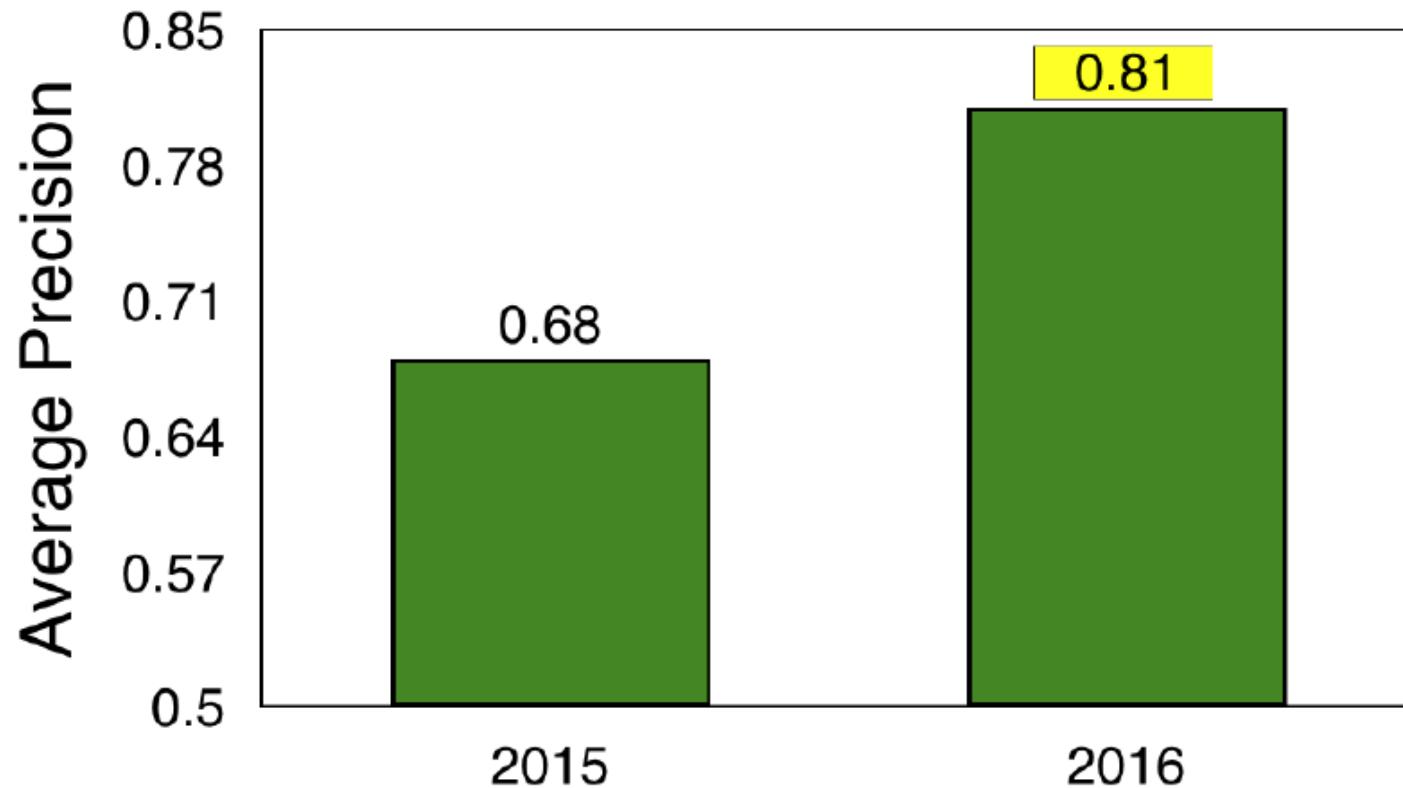
# Detection



Fully annotated 30 object classes across 6,278 snippets (train+test)

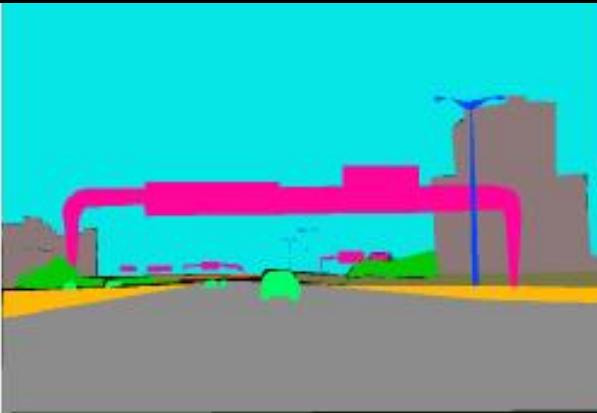


## Video Detection



# 图像识别 | Beyond ImageNet

# Sense Classification & Parsing



# Image Captioning

A person riding a motorcycle on a dirt road.



Two dogs play in the grass.



人脸识别 | 进入混战模式

视频分析 | 门槛在慢慢降低

视频预测 | 效果还比较诡异

基础 | 前沿 | 应用 | 评价



SECURITY  
CHINA

# 2016 Security China

## 中国国际社会公共安全产品博览会

2016 China International Exhibition on Public Safety and Security

中文版 | English | Русский



OK



服务经济民生

展示安防技术

构建智慧城市

推进平安中国

网站首页

关于展会

展会活动

展会新闻

参展服务

观众服务

交通及住宿

下载中心

联系我们

主办单位：中国安全防范产品行业协会

展览日期：2016年10月25-28日

展览地点：中国国际展览中心(新馆)

# 旷视智能安防解决方案

Face++ 旷视



# 智造平安城市！ 旷视智能安防解决方案——天眼系统

Face++ 旷视



**SENSEFACE**

抓拍人数 1009

告警人数 46

视频路数 4



人脸检测

人脸检测

属性

风格



定位：展会1号相机  
10-25-2016



任务：人脸分析系统1  
11:05



男  
19岁  
黄种人  
睁眼  
闭嘴



男  
27岁  
眼镜  
黄种人  
睁眼  
闭嘴



男  
25岁  
眼镜  
黄种人  
睁眼  
闭嘴



定位：展会1号相机

任务：人脸分析系统1



定位：展会2号相机

任务：人脸分析系统2



定位：展会3号相机

任务：人脸分析系统3



定位：展会4号相机

任务：人脸分析



展会：2号相机 2016-10-



姓名：王晶  
性别：男  
所在库：安博展员工库

展会：3号相机 2016-10-



姓名：梁博  
性别：男  
所在库：安博展员工库

**SENSEFACE**

抓拍人数 1031 告警人数 48

属性 风格

10月25日 2016 11:06:11

任超 人脸识别系统1

性别 年龄 种族 表情 状态

男 26岁 黄种人 微笑 闭眼

女 23岁 男 19岁 黄种人 微笑 闭眼

男 23岁 男 22岁 黄种人 微笑 闭眼

男 23岁 喜悦 黄种人 微笑 闭眼

姓名 梁博 性别 男 所在库 安博展员工库

展会1号相机 2016-10-25 11:06:11

展会2号相机 2016-10-25 11:06:11

展会3号相机 2016-10-25 11:06:11

展会4号相机 2016-10-25 11:06:11

视频路数 4

2016-10-25

# SenseFace-TX1

## 硬件

NVIDIA Tegra X1

核心板仅一张信用卡大小

质量约75g

功耗6.5 ~ 15W

## 功能

多人脸检测

人脸关键点提取

人脸质量评估

动态跟踪

## 性能

1080P监控视频

25hz实时处理



```
./cs_senseface.sh: line 2: 16755 Aborted
sensetime@sensetime-Super-Server:~/Desktop$ ./cs_senseface.sh
(release version
-----query db task: i=0, taskName=人脸分析系统1, vsName=展会1号相机, type1, path1:tcp://admin:admin123@192.168.0.200, state1:open, rssi1:7.7, desc1:未知
-----query db task: i=1, taskName=人脸分析系统2, vsName=展会2号相机, type1, path1:tcp://admin:admin123@192.168.0.200, state1:open, rssi1:7.7, desc1:未知
-----query db task: i=2, taskName=人脸分析系统3, vsName=展会3号相机, type1, path1:tcp://admin:admin123@192.168.0.200, state1:open, rssi1:7.7, desc1:未知
-----query db task: i=3, taskName=人脸分析系统4, vsName=展会4号相机, type1, path1:tcp://admin:admin123@192.168.0.200, state1:open, rssi1:7.7, desc1:未知
libpng warning: iccp: known incorrect sRGB profile
E1025 11:21:40.573937 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.582726 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.590593 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.590679 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.699939 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.707291 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.720096 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.720065 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.736455 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.743432 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:40.754927 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.093533 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.093899 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.101660 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.108548 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.162178 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.171679 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.177114 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.188246 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.195377 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.201452 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
E1025 11:21:41.208427 16821 senseface_app.cpp:164] [VerifyClient] dbName: 222222296f0bc3b0f7942578000471ack:b221
OpenCV Error: Assertion failed (0 <= rot.x && 0 <= rot.y && 0 <= rot.z && n.cols && 0 <= rot.y && 0 <= rot.height && rot.height = n.rows) in file /home/appuser/software/opencv-2.4.10/modules/imgproc/src/matrix.cpp, line 323
terminate called after throwing an instance of 'cv::Exception'
what(): /home/appuser/software/opencv-2.4.10/modules/imgproc/src/matrix.cpp:323: error: (-215) 0 <= rot.x && 0 <= rot.y && 0 <= rot.z && n.cols && 0 <= rot.y && 0 <= rot.height && rot.height = n.rows
n.rows in function Mat
```



基础 | 前沿 | 应用 | 评价

评价 | 泡泡越大，失望越大 (AGI)

评价 | 能解决的问题是有限的

评价 | 自动编程，呵呵哒

# 评价 | 警惕伪科学和各类歧视

评价 | 新交互开启新攻击面

谢谢各位抽出时间来一起讨论

有问题可以提问交流

# 深度学习：基础，前沿，应用，评价

吴伟

lazyparser@gmail.com  
中国科学院软件研究所