

Code Speed Tracking Infrastructure for compilers, runtimes, libs

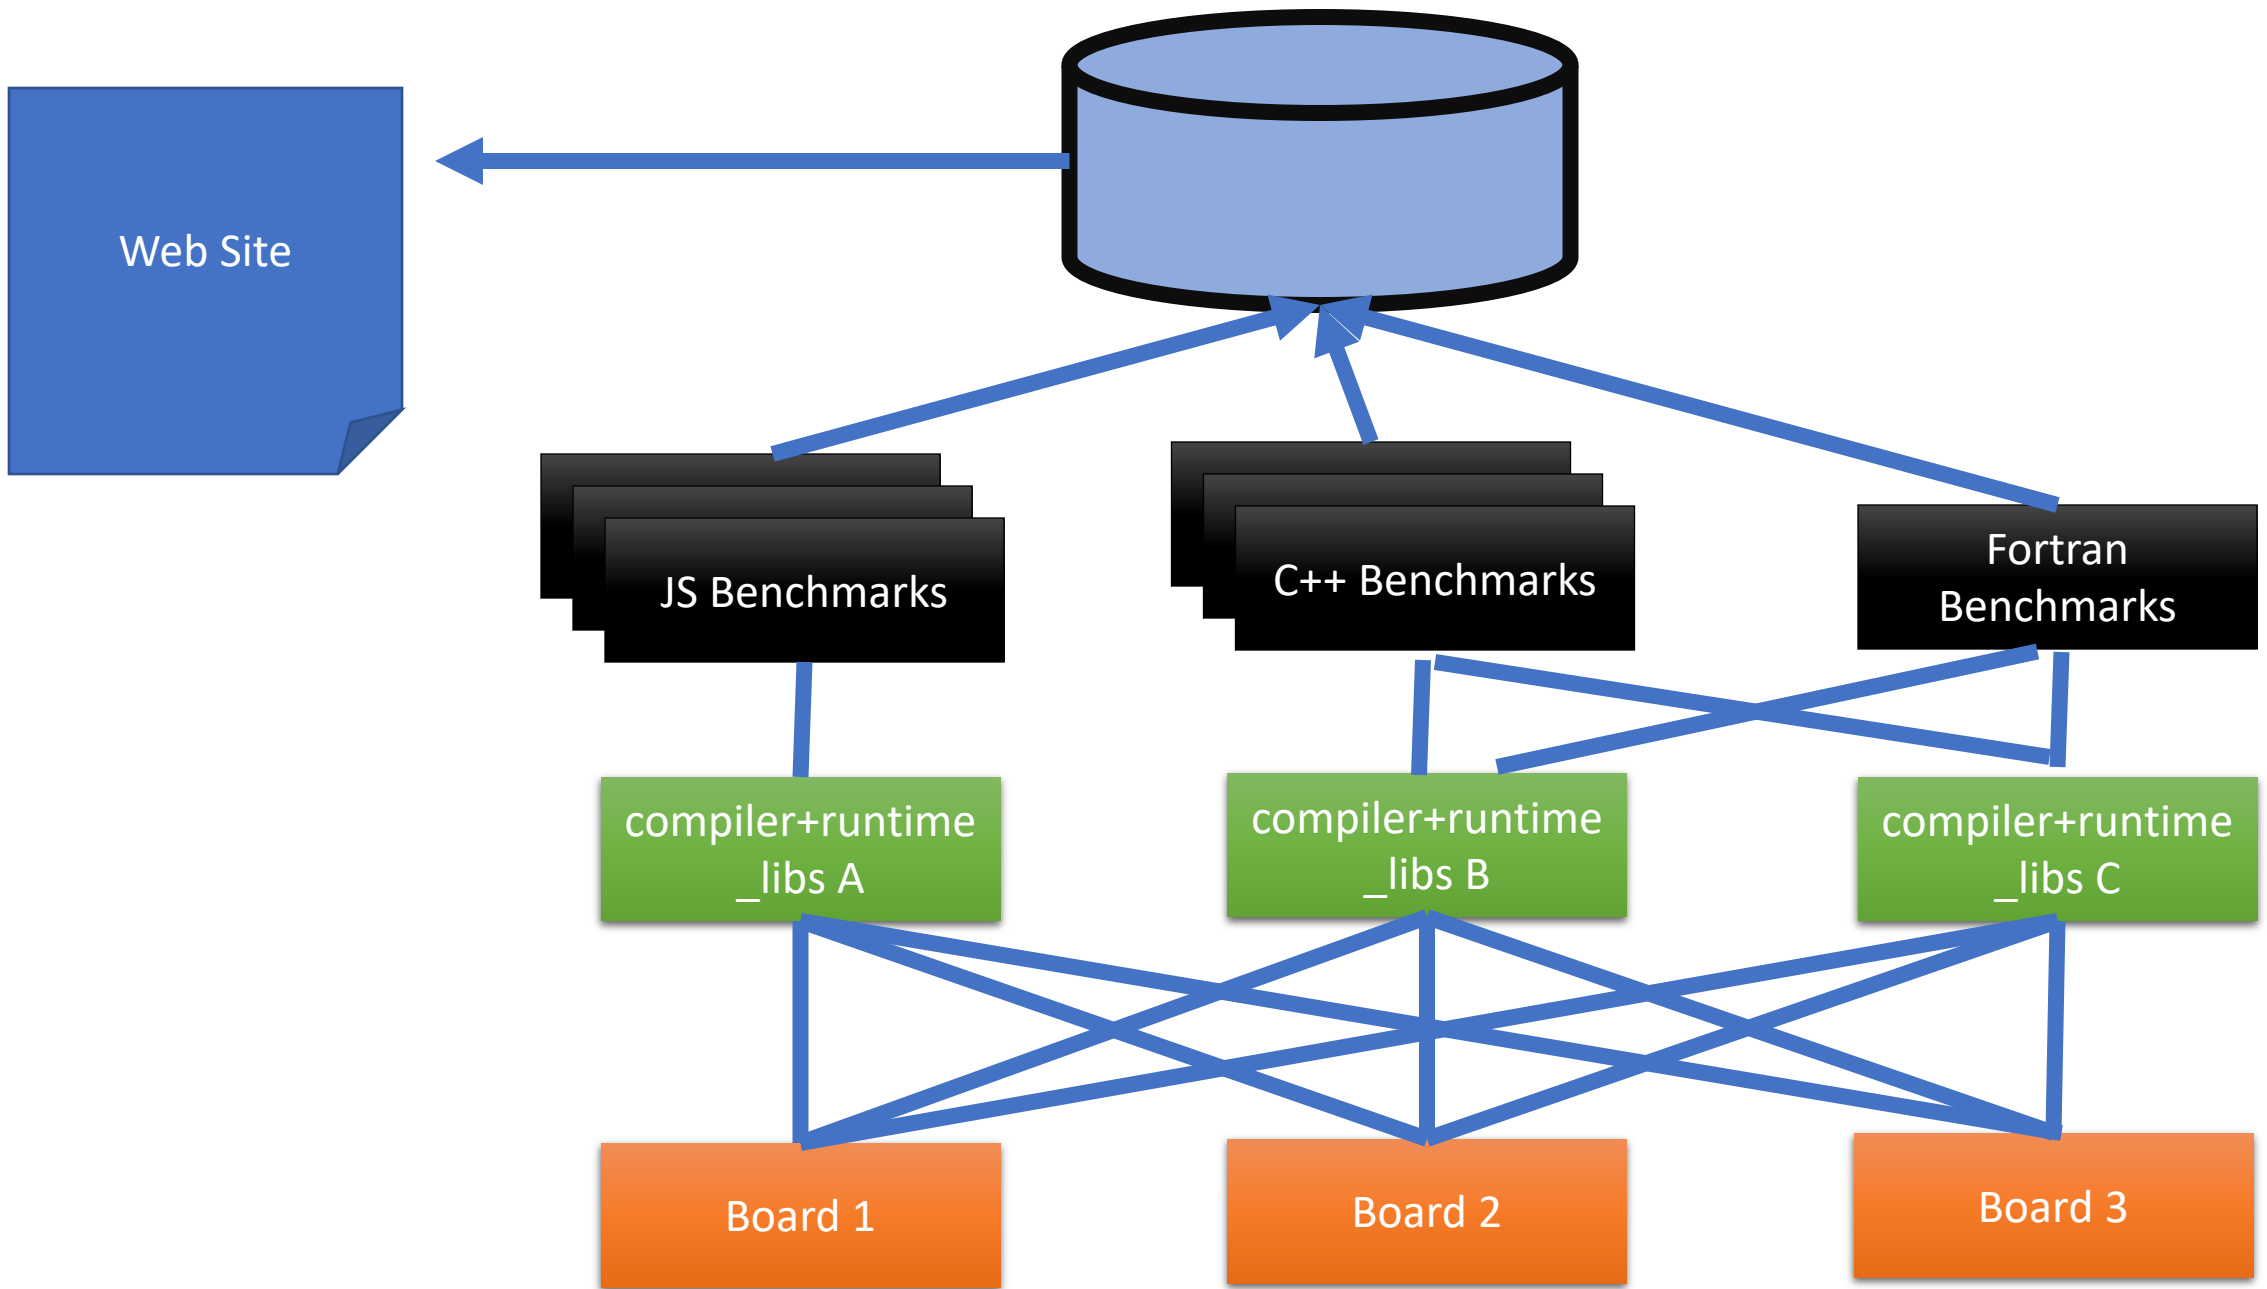
RFC Proposal, first version

Wei Wu lazyparser@gmail.com

2020-12-15

Overview - scope

- Build a performance tracking system for RISC-V ecosystem
 - Running on physical boards. Get real performance data.
 - Hifive Unleashed is ready. Hifive Unmatched and PolarFire icicle will follow.
 - Targeting GCC, Clang/LLVM, V8, Spidermonkey, OpenJDK, Rust, Golang, and LuaJIT, etc.
 - Also targeting OpenBLAS and other important libraries in future
 - Using Both free and commercial Benchmarks
 - Free: mBench? SunSpider, Octane, Kraken
 - Non-free: SPEC CPU 2006, etc.
 - A website to show the data points. Using arewefastyet open-source project



Demo - <https://arewefastyet.com/linux64/benchmarks/octane?numDays=60>

Platform

Linux 64bit ▾

Category

Benchmarks ▾


Results

Octane (JS shell) ▾


Time range

Last 14 days ▾

Octane (JS shell) [↗](#) [↓](#)



octane-Box2D [↗](#)



RISC-V market justification

- The Code Speed (performance) is blocking the prosperity of RISC-V
 - Java speed on RV was ~100x slower than x86/Arm w/o JIT
 - V8 and other runtimes w/ JITs are under optimized for RISC-V
 - GCC and Clang still have huge potential to improve
 - So, code speed is important.
- A platform tracking all free available toolchains & runtimes & libs
 - It is important to know where we are, how far can we achieve.
 - It is kind of public goods, lighthouse for toolchains & runtimes & libs
 - So, it is better to put under the RISC-V International

Phasing

- Phase 1: set up the infrastructure
 - one compiler (v8), one board (unleashed), three (JS) benchmarks
 - One engineer, in one month. Run the process. Make the website online.
- Phase 2: more toolchains & runtimes, all available open-source benchmarks.
 - Include upstream GCC, RISC-V GCC, upstream Clang/LLVM
 - Include OpenJDK/HotSpot, OpenJDK/OpenJ9
 - Include Rust
 - Include Go
 - Include mBench and all available open-source benchmarks

Phasing - cont.

- Phase 3: more physical boards and libraries
 - Add Hifive Unmatched and other physical boards
 - Set up APIs or processes to let RVI members run the scripts on their own devices and upload to the platform/website.
 - Include OpenBLAS and other compute libs.

Outline plan and timescale

- Phase 1: 1 engineer. 1 month. Plan: 2021-2-1 ~ 2021-3-1
 - Leave one month for volunteers. Can start today, although.
- Phase 2: 2 engineer. 3 month. Plan: 2021-3-1 ~ 2021-6-1
- Phase 3: 2 engineer. 2 month[1]. Plan: 2021-6-1 ~ 2021-8-1
 - [1] it depends on the status of hardware.

Prerequisites

- Basic knowledge about performance testing and tuning
- Basic knowledge about how to establish a website and tuning JS and PHP code.
- Have RISC-V physical board (not necessary if you just want a test)
- Basic knowledge about how to write bash and python scripts.

Deliverables

- An open-source toolset for building compilers & running benchmarks & show perf results on web.
- Data reports generated by this tracking platform.
- A monthly report to the SIG is a routine deliverables for all projects, which will cover
 - key progress in the past month
 - table showing tasks to be done, in progress, complete and verified
 - test/benchmark results
 - updates to the risk register
 - plans for the next month

Milestones

- I think it is just a small project in the first place so the milestone is the phase stage itself.

Costs

- equipment costs
 - Can we just get some boards donated by the member companies made them?
 - An VPS/laaS for hosting the website. PLCT has existing VPS and can donate for it.
- license costs
 - Use free & open benchmark first.
 - How to donate? Is it possible? (due the LINCENSE of each benchmarks)
- personnel costs
 - TBD. Roughly 2 person 1 day per week?

Risk register

- TBD.
- Risks are assessed by the Impact (I) they have on the project from 1 (minor) to project killer (3) and by the Likelihood (L) of the risk occurring from 1 (10% chance) through 10 (100% chance). The two are multiplied to give an overall Risk Factor (R). Mitigation must be provided for any risk with $I = 3$ or $R \geq 10$.

Support

- Physical boards are needed.
 - Especially the boards that can run Linux are welcome.
- Members can run scripts in their own boards and upload the data.
- Need commercial toolchains & runtimes run the scripts and send back the performance data to the tracking platform.
- A project is only useful if members are willing to back it with real resources, which can be in kind or in cash. Generally a project should only go ahead under the aegis of RISC-V if multiple members are backing it.
- Table of member organizations and commitments

Welcome to contribute

- This is the first-time proposal.
- It is free to grab. If anyone want to do it, can just rise your hand and bring a detailed proposal next time.
- Need to create a TG? (needs SSC's approval)